



Deep Learning with Python and Friends

Douglas Starnes

Indy.Code() 2019

Polyglot ninja

Memphis, TN area

Co-director of MemPy

Data Science/Machine Learning, Mobile Apps,
wannabe Game Designer

Pluralsight Author

Watch for carefully placed marketing through the talk!

@poweredbyaltnet



@poweredbyaltnet

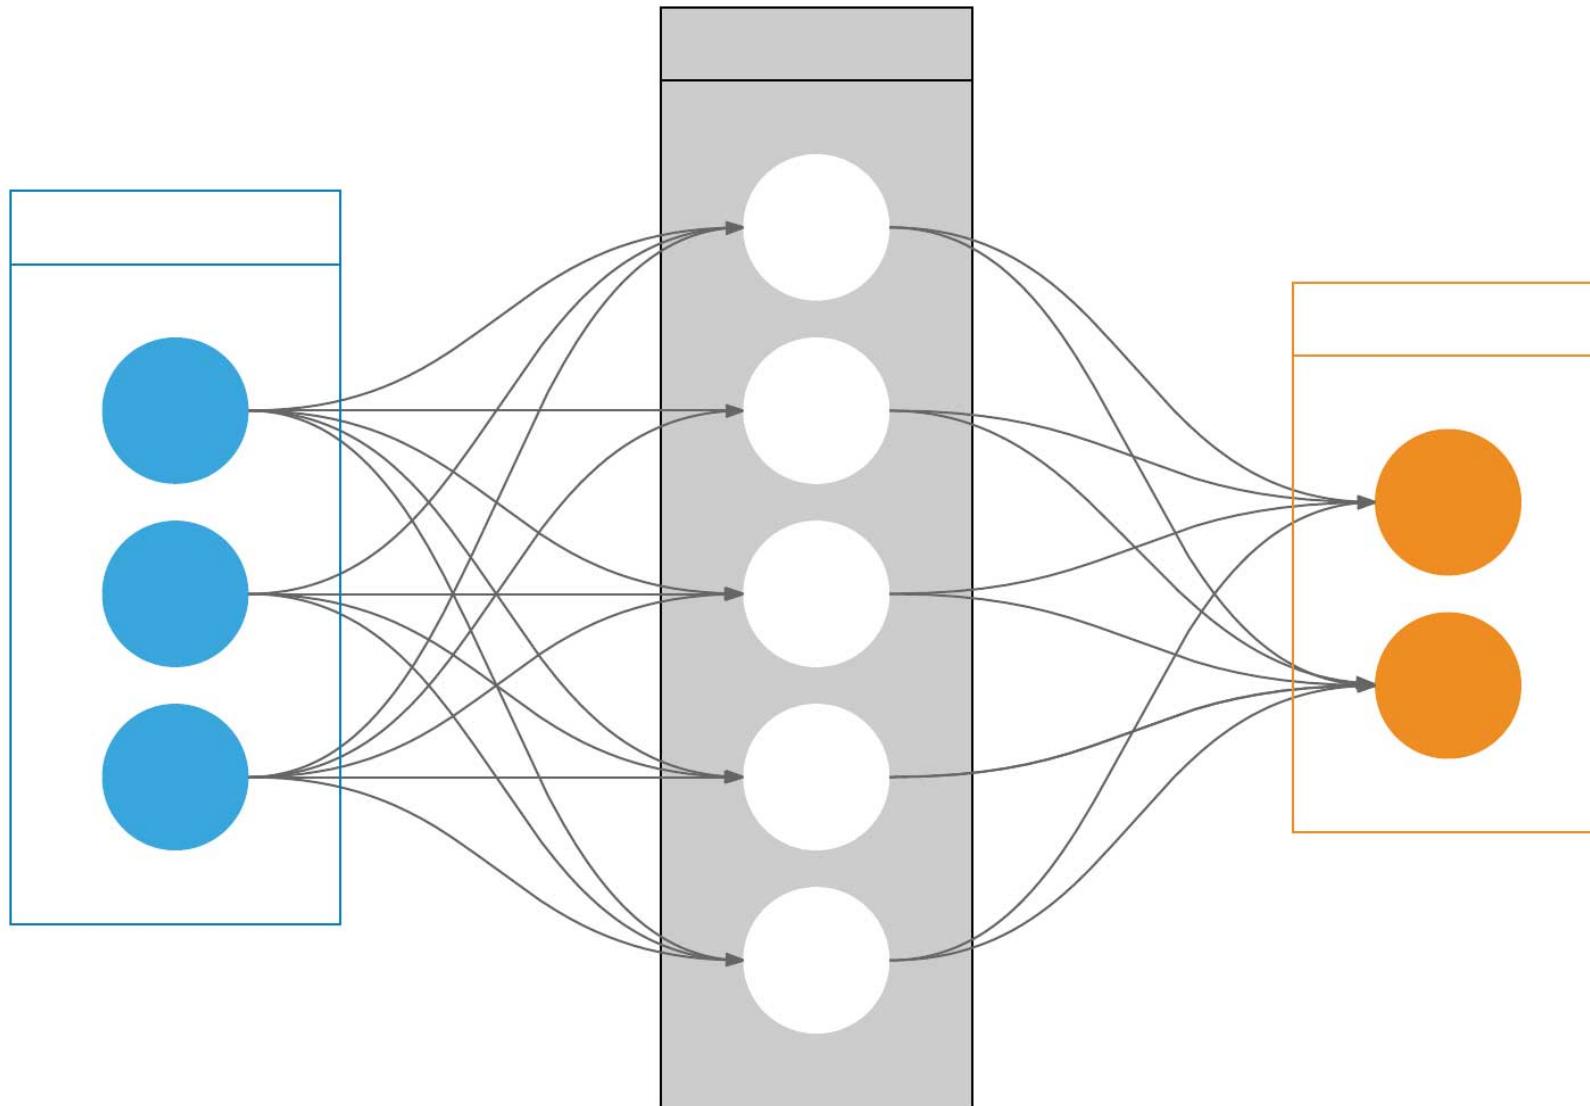
Artificial Intelligence

Machine Learning

Deep Learning

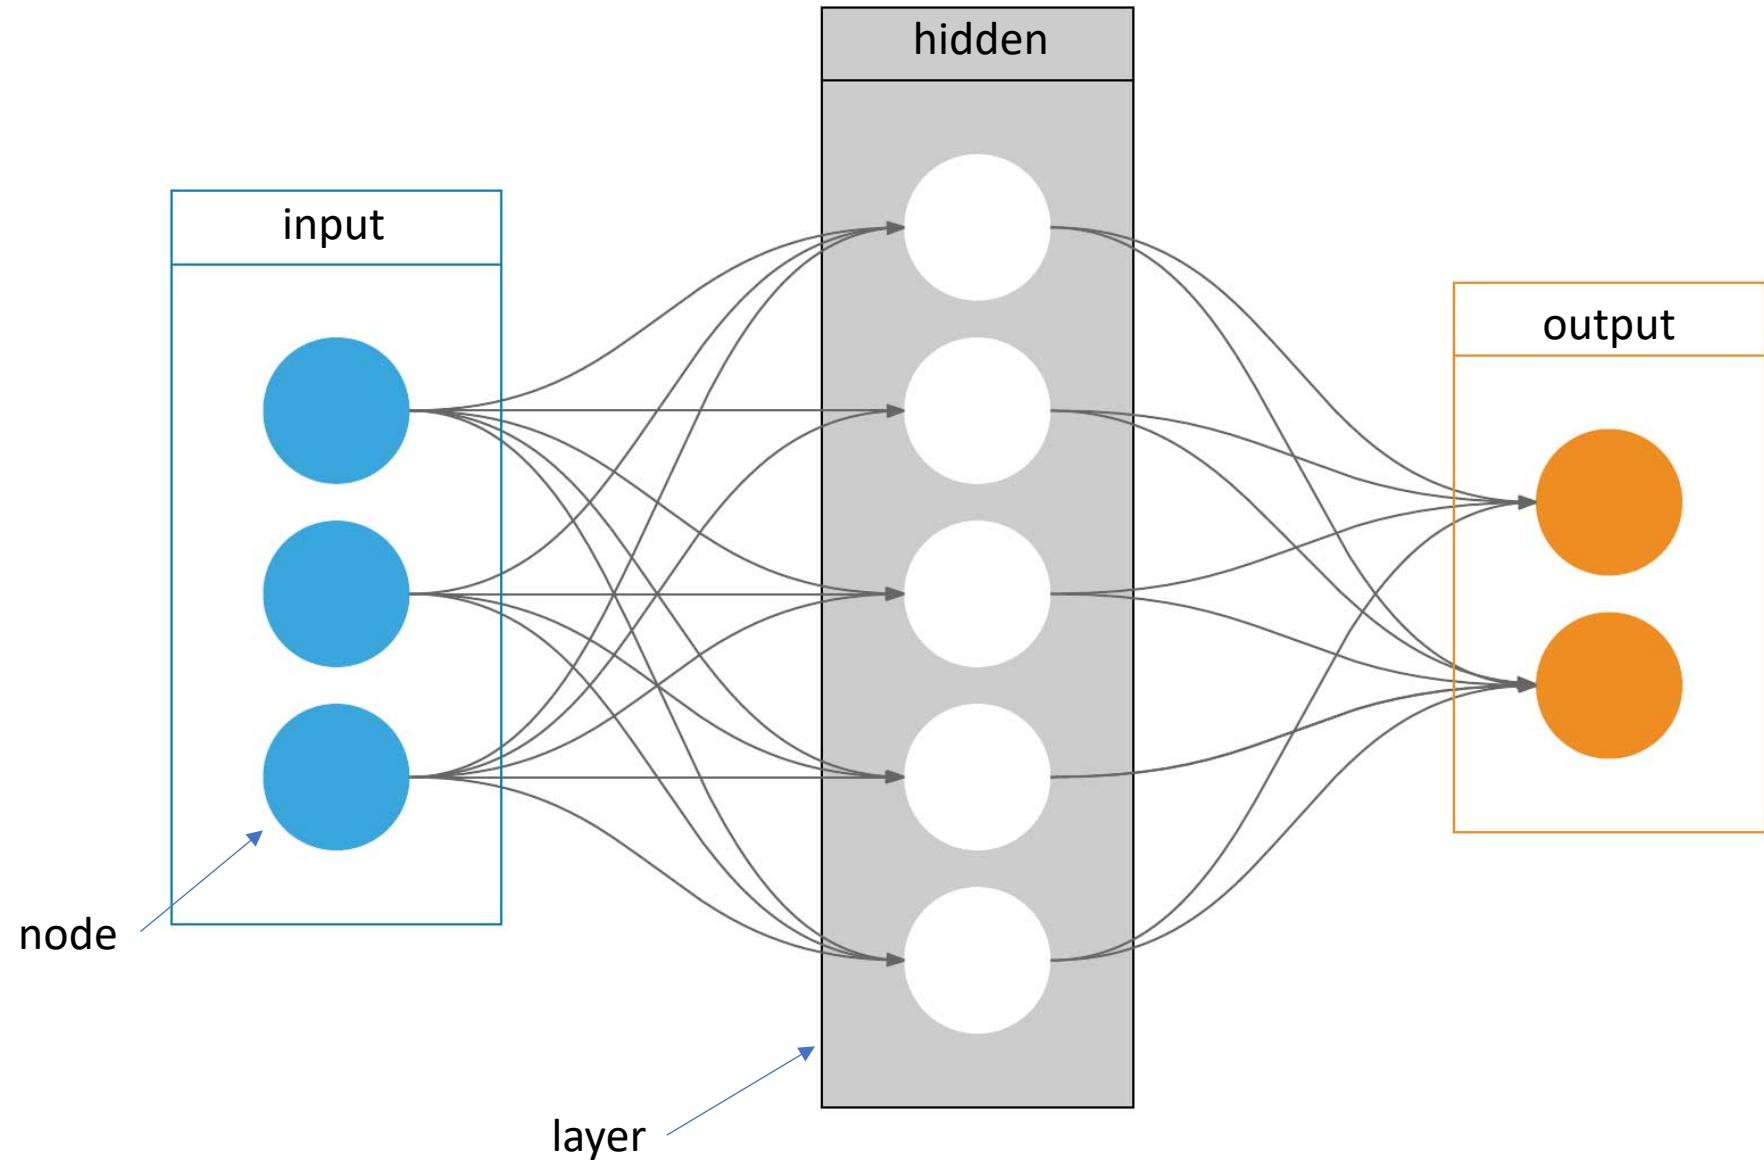
D
a
t
a

S
c
i
e
n
c
e

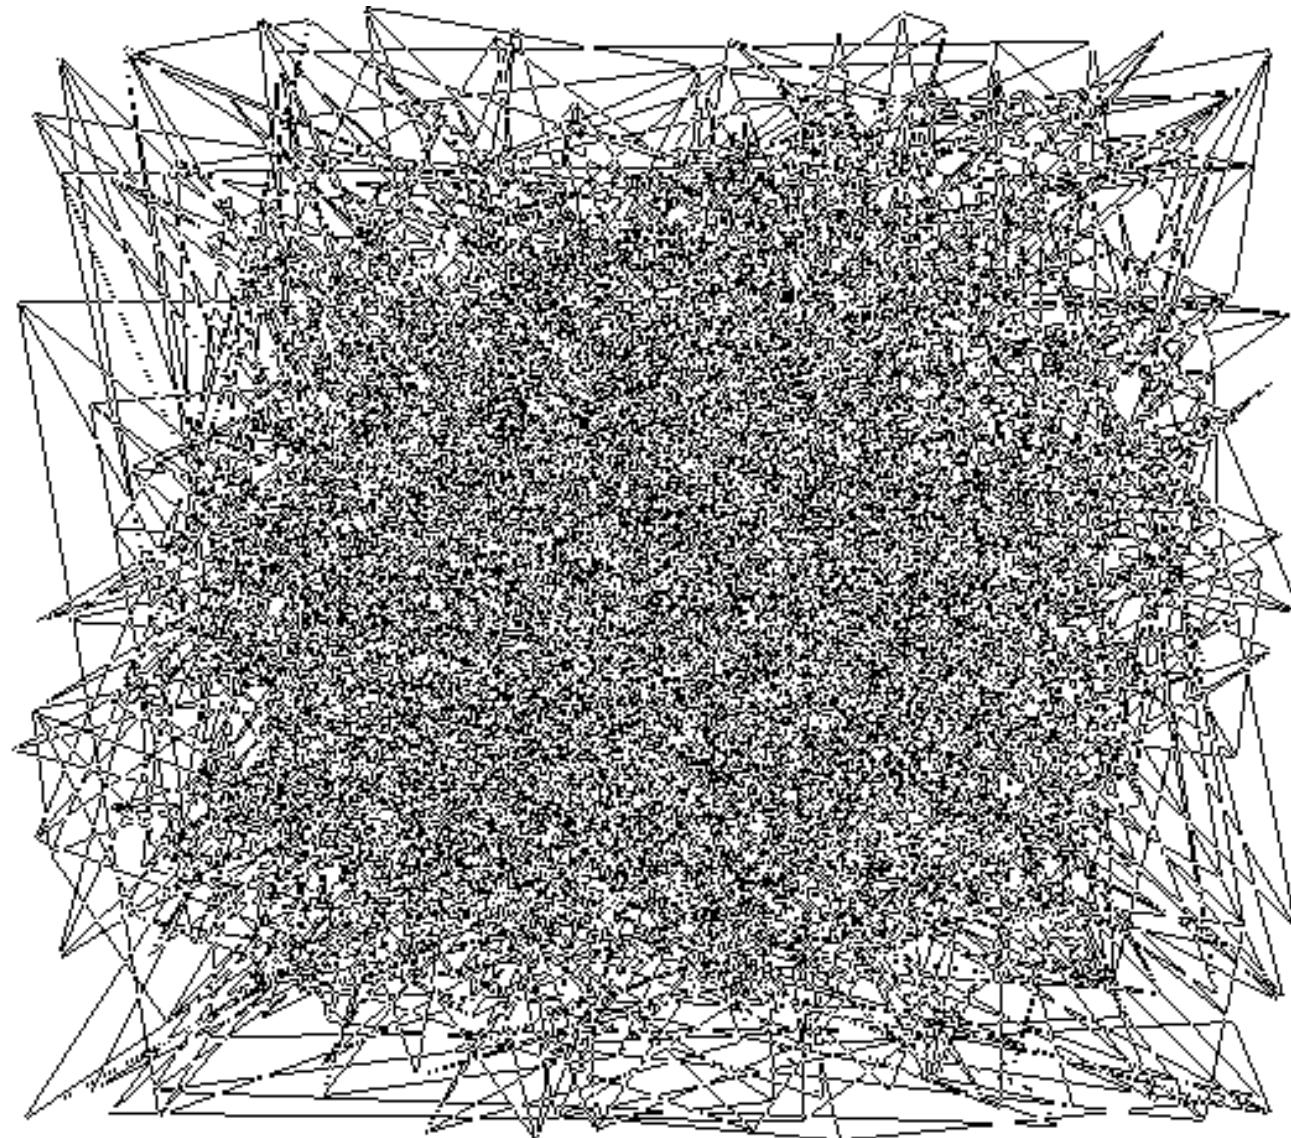


Neural networks are loosely inspired by how neuroscientists today think the brain might possibly work.
(assuming they are on the right track)





A neural network with more than two hidden layers is called a **deep** neural network.



feature vector				target
Sepal Length	Sepal Width	Petal Length	Petal Width	Species
5.1	3.5	1.4	0.2	setosa
...
7.0	3.2	4.7	1.4	versicolor
...
5.7	2.8	4.1	1.3	virginica

rows -> observations

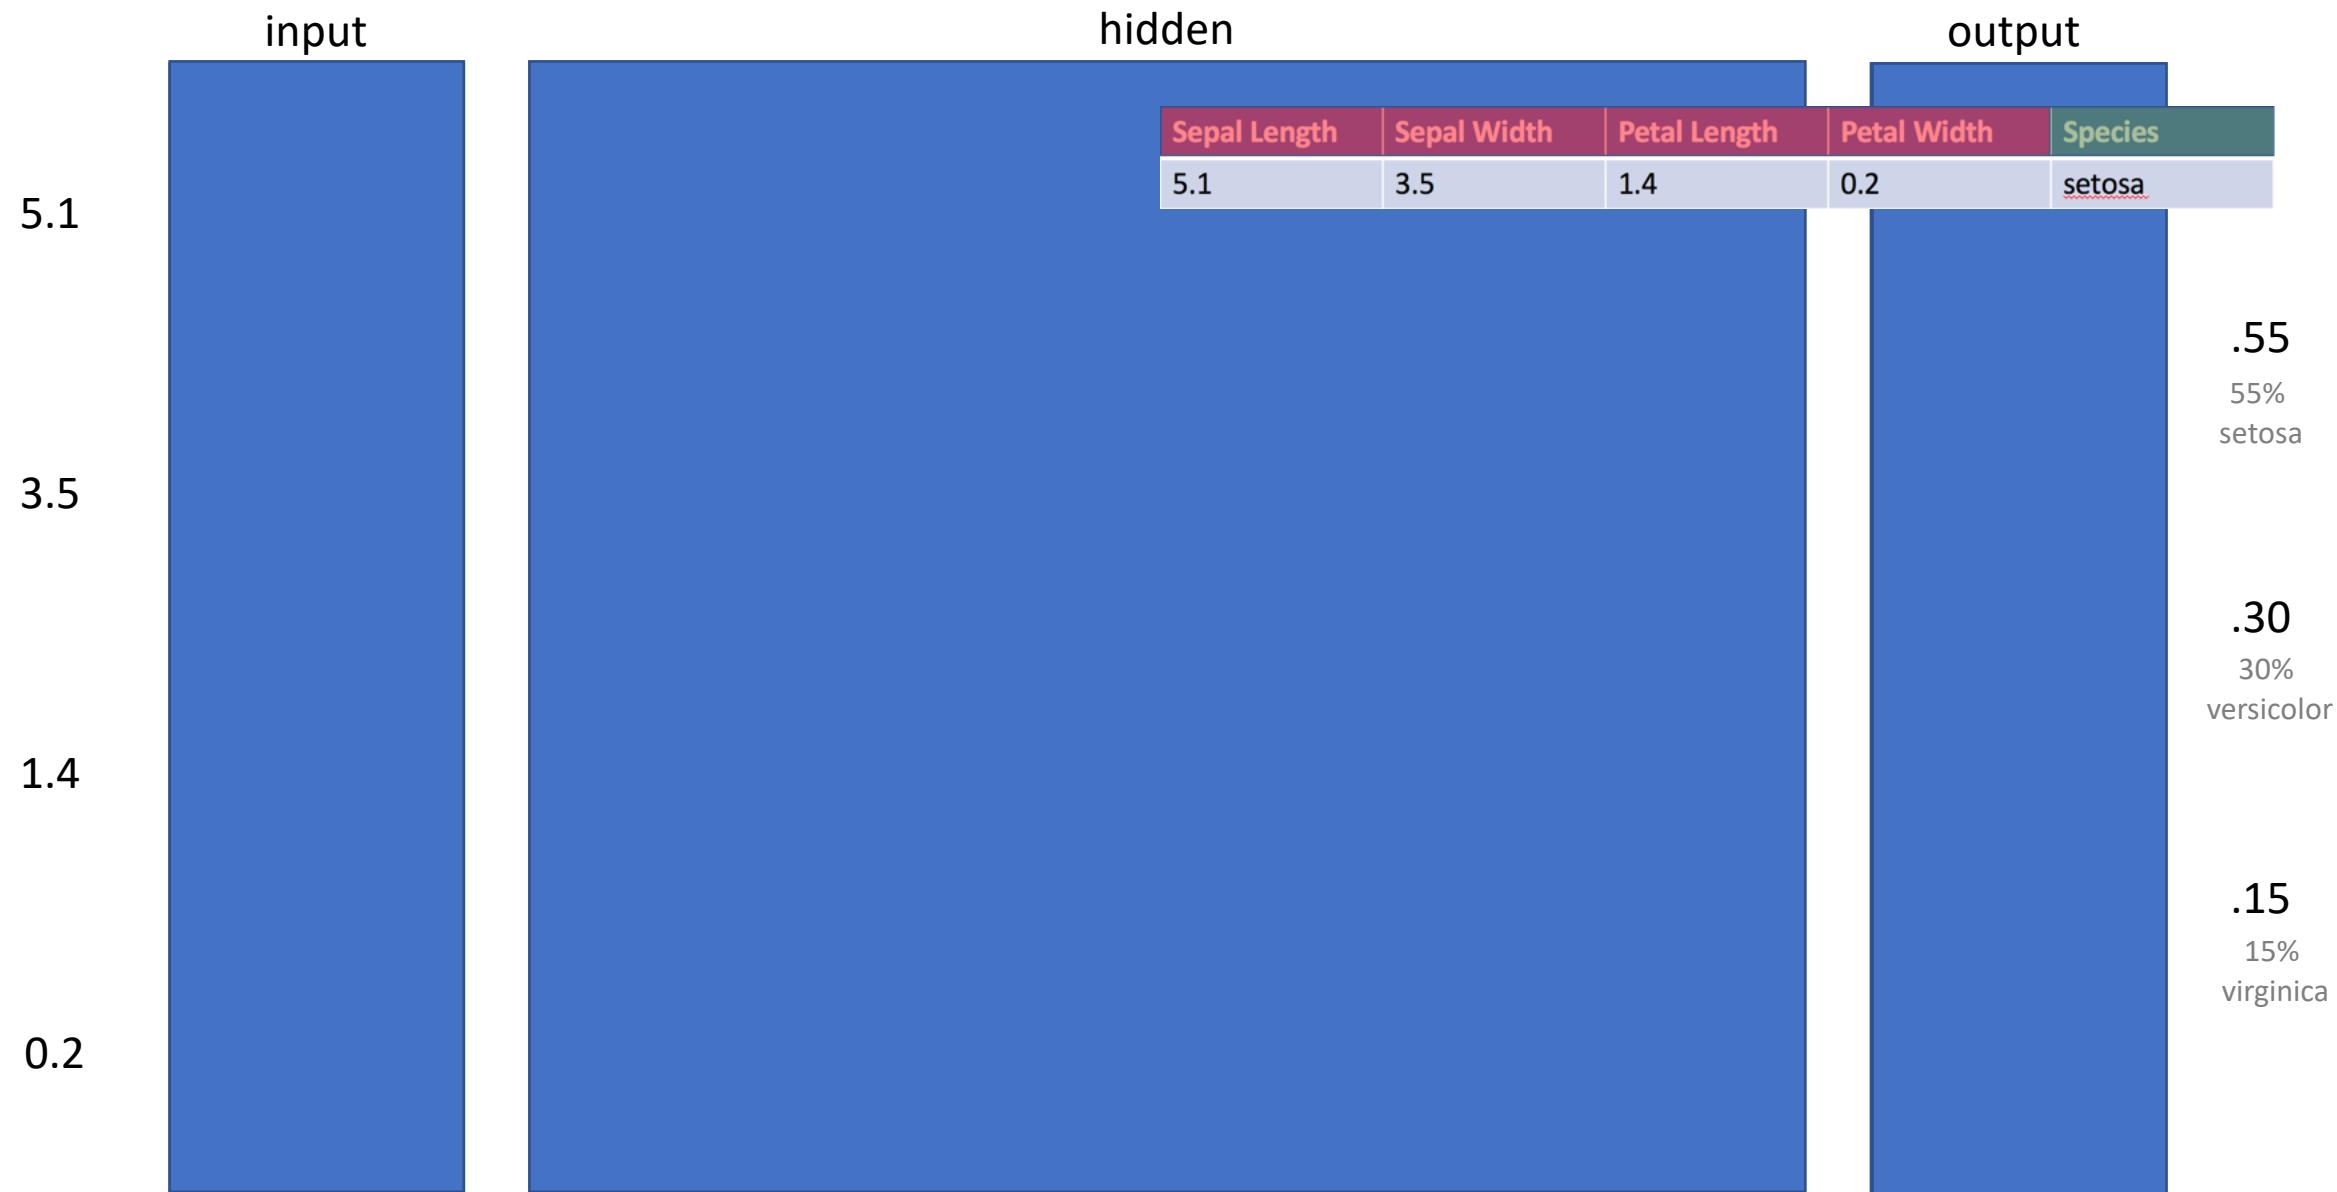
columns -> features

feature vector / independent variables -> observed values

target / dependent variable(s) -> what we want to predict

train a **model** that, given a **feature vector**, will (hopefully accurately) assign a **target**

train a **model** that, given values for *Sepal Length*, *Sepal Width*, *Petal Length*, *Petal Width*, will assign one value from the set of *setosa*, *versicolor*, or *virginica*





Is this
prediction
correct?



WHO
CARES

3

90 00 15 00 15

There is only one trick question in
the entire talk. This is it.



During training, we are not concerned with making a correct prediction, but rather making a prediction that is as close to correct as possible.



WTF?

@poweredbyalt.net

Prediction

Actual

.55

1.0

.30

0

.15

0



I think we
can do
better!

Prediction

.9999999

.0000001

.0000001

Actual

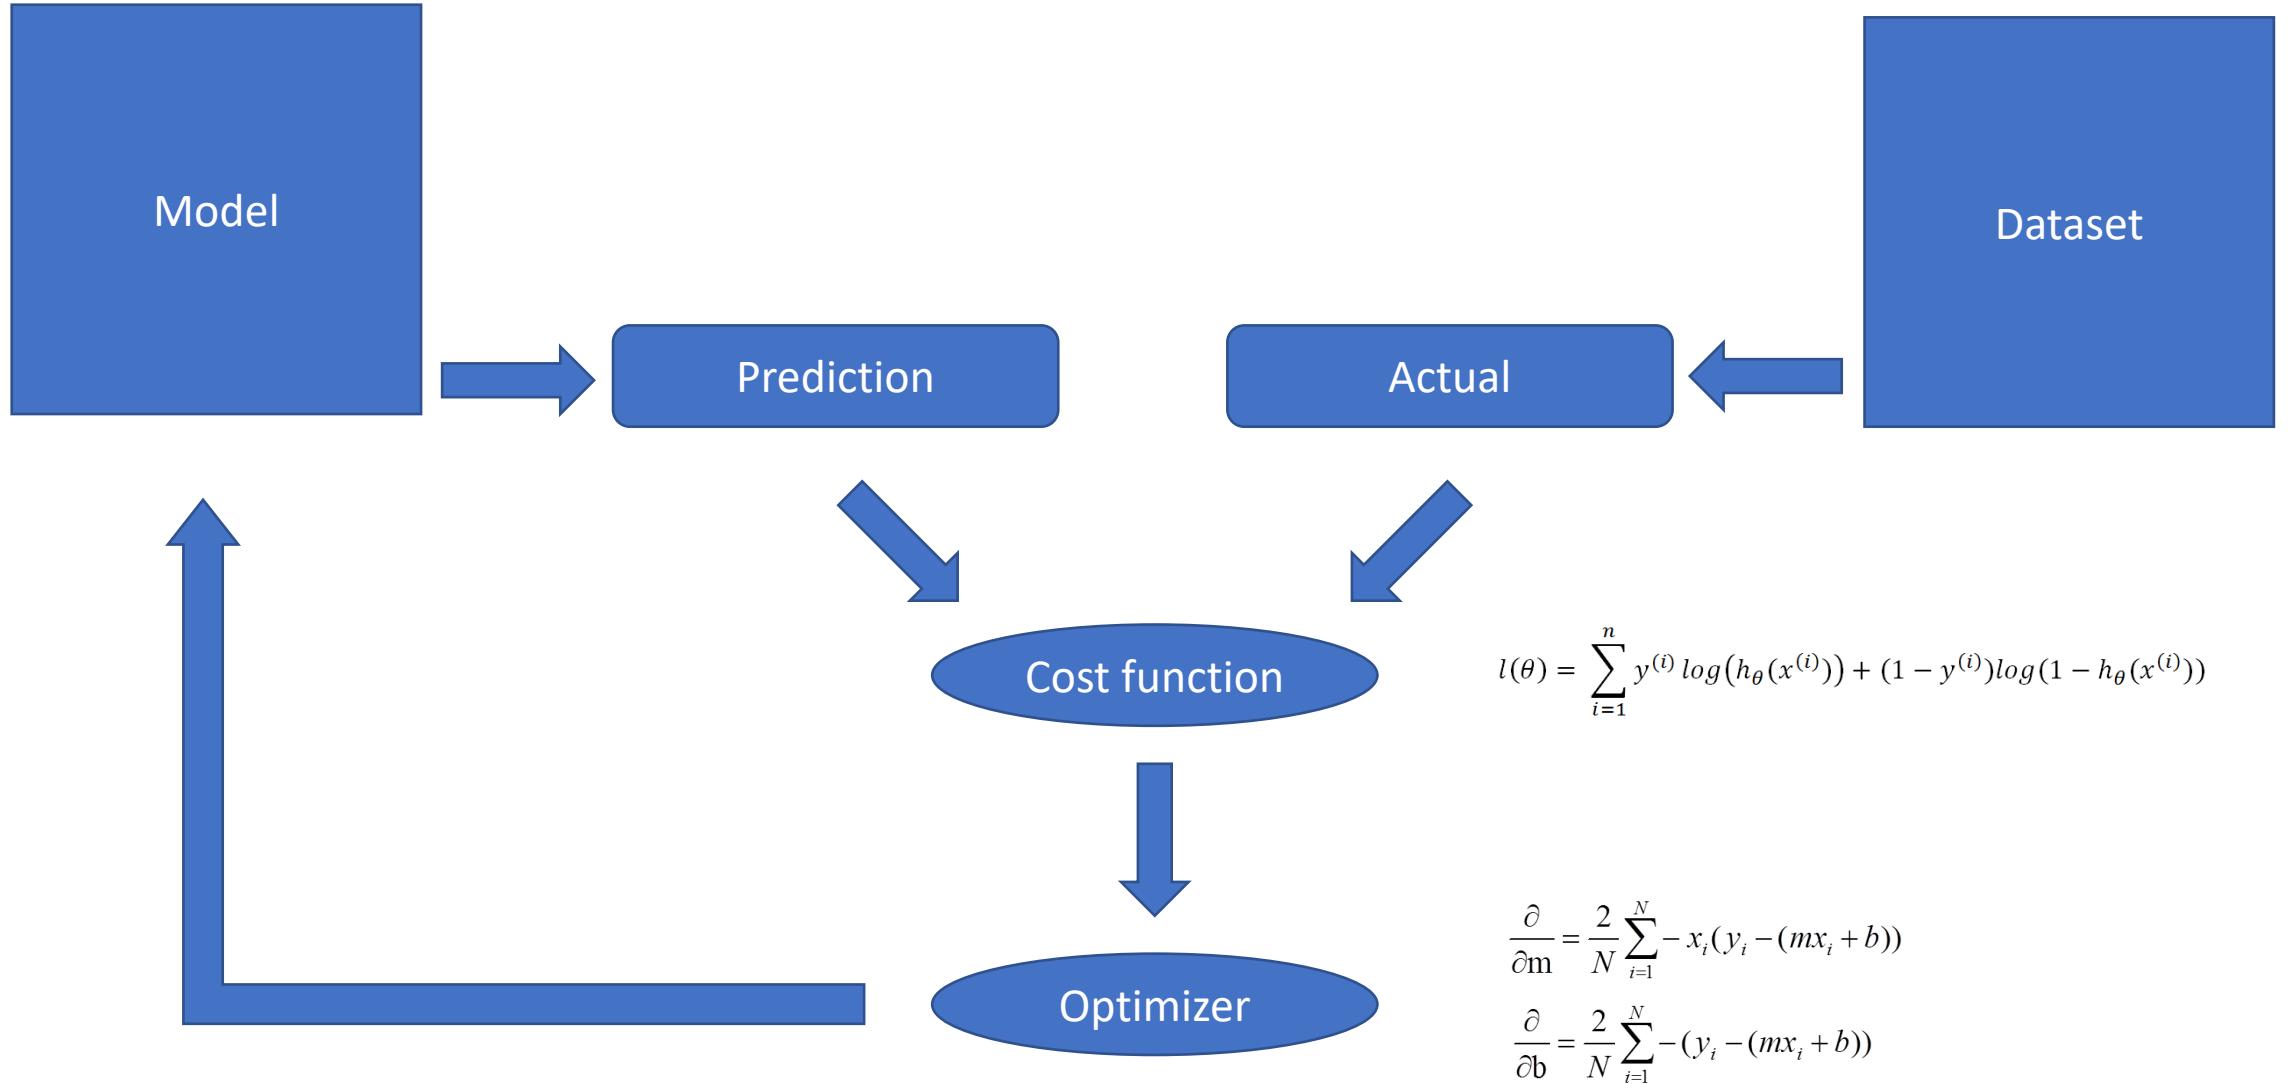
1.0

0

0



More 9's please!

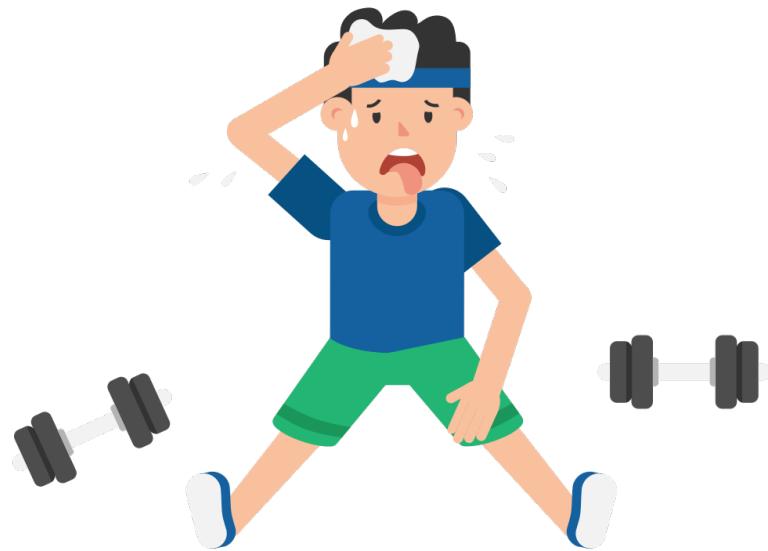


$$l(\theta) = \sum_{i=1}^n y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))$$

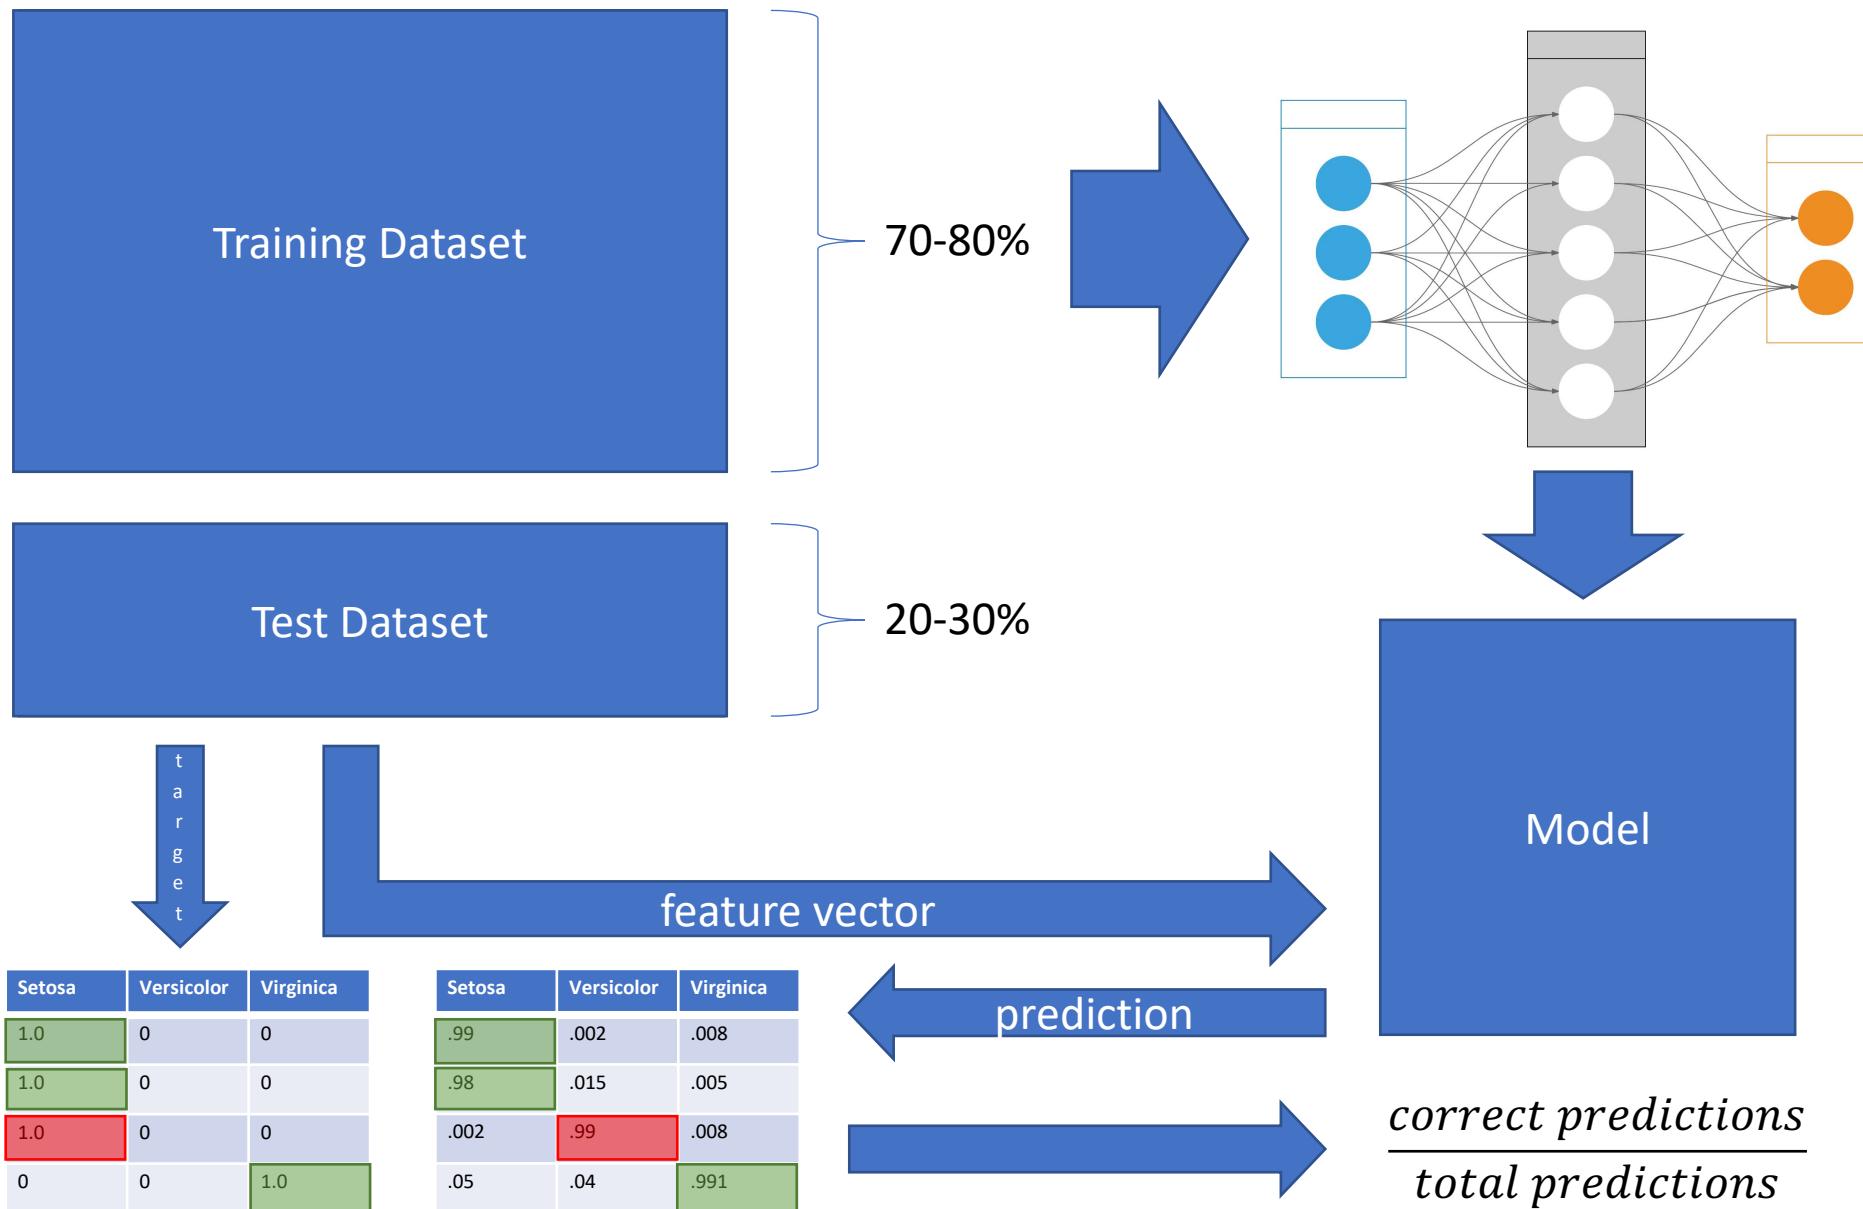
$$\frac{\partial}{\partial m} = \frac{2}{N} \sum_{i=1}^N -x_i(y_i - (mx_i + b))$$

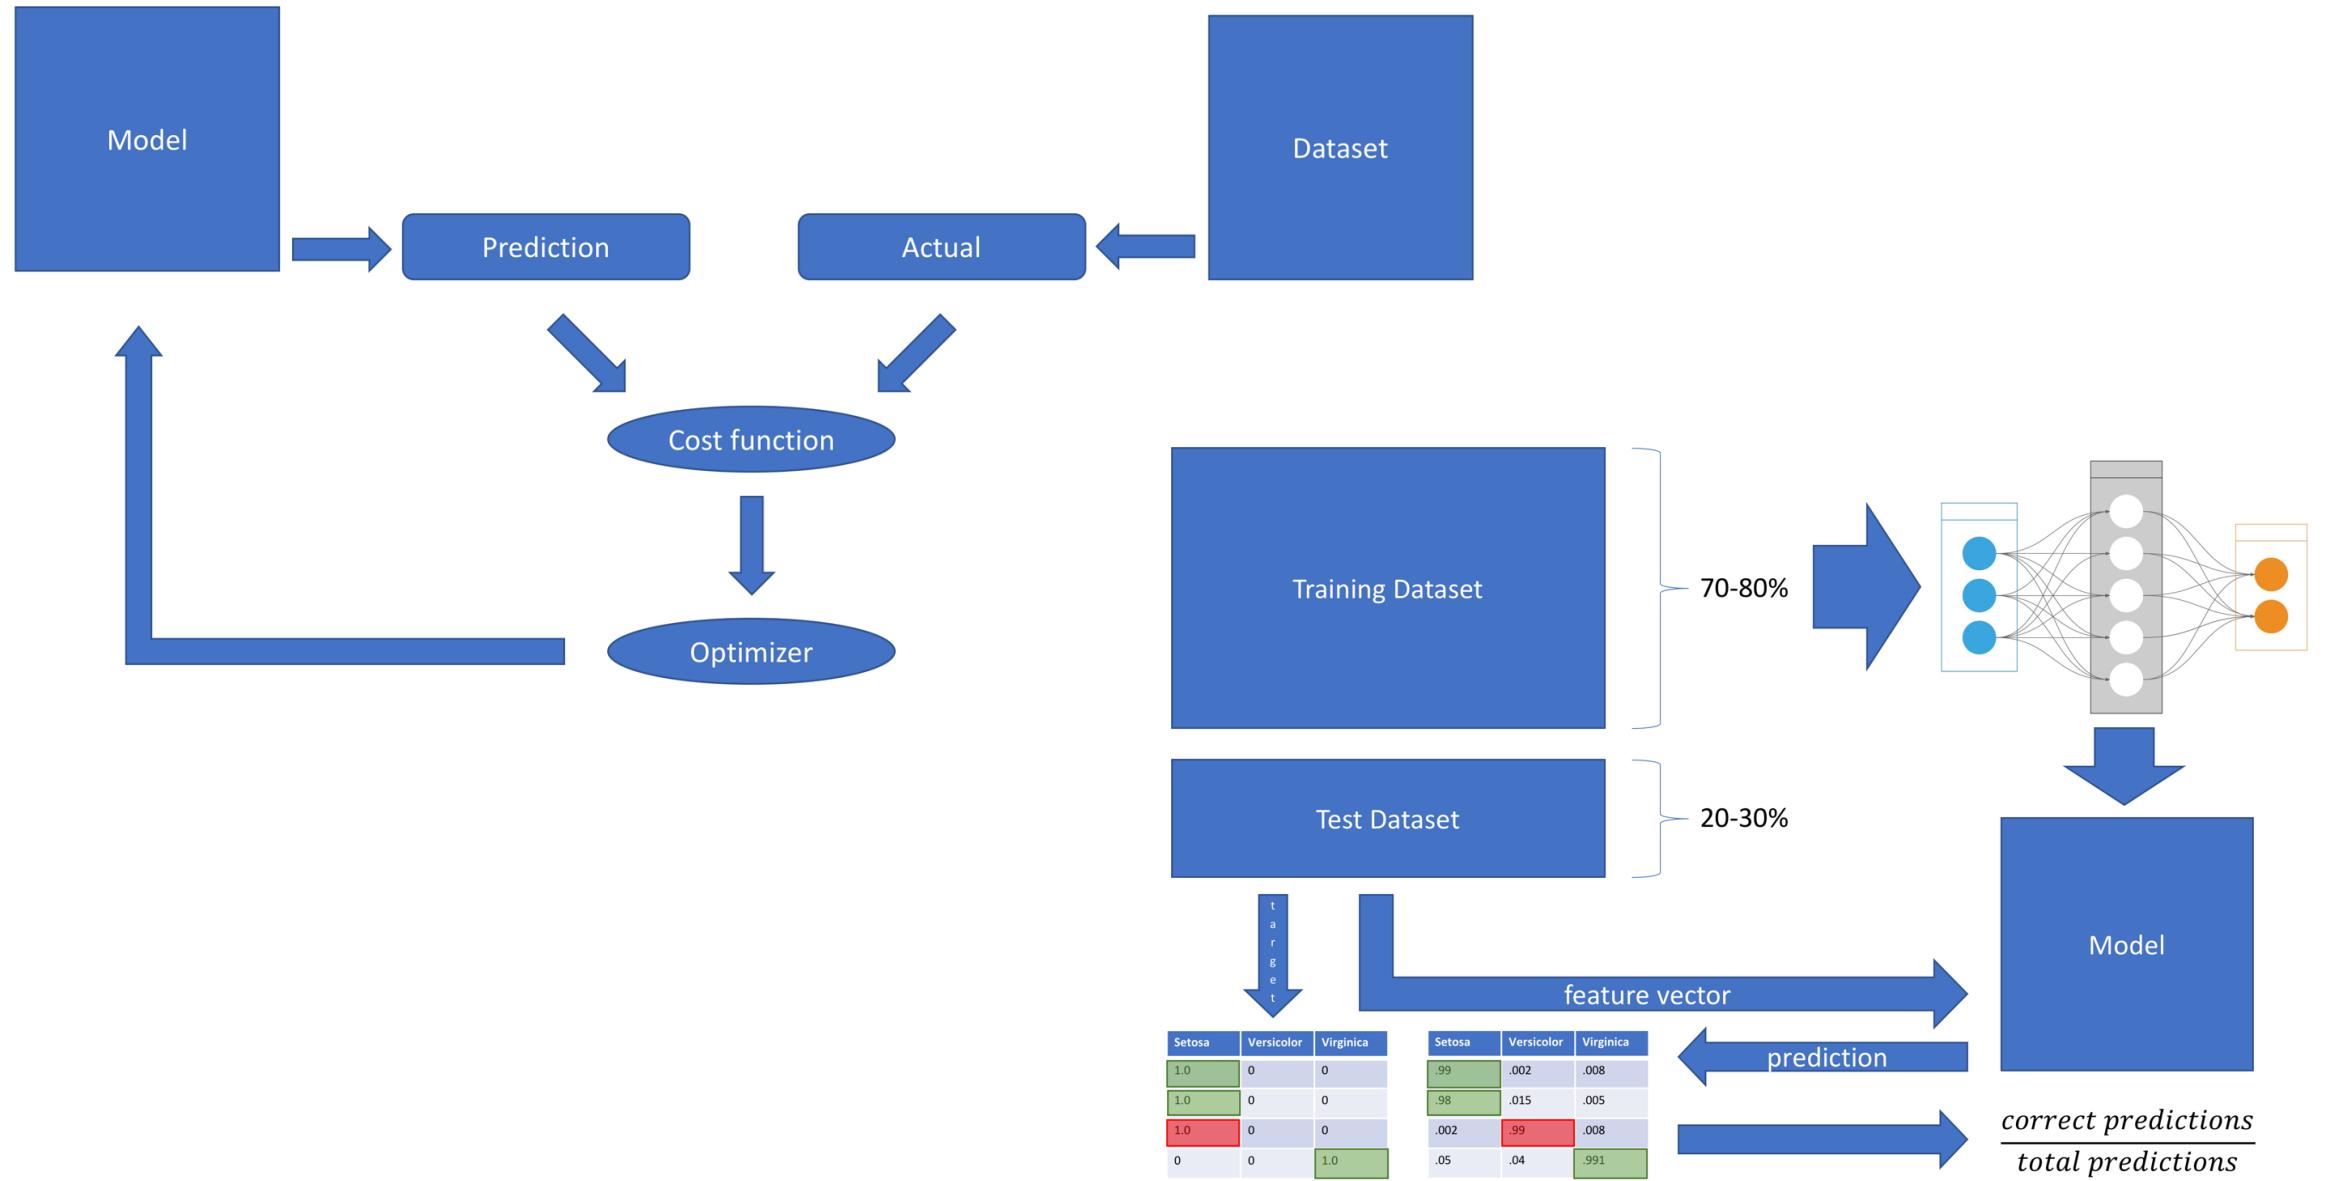
$$\frac{\partial}{\partial b} = \frac{2}{N} \sum_{i=1}^N -(y_i - (mx_i + b))$$

The model is trained



Now what?

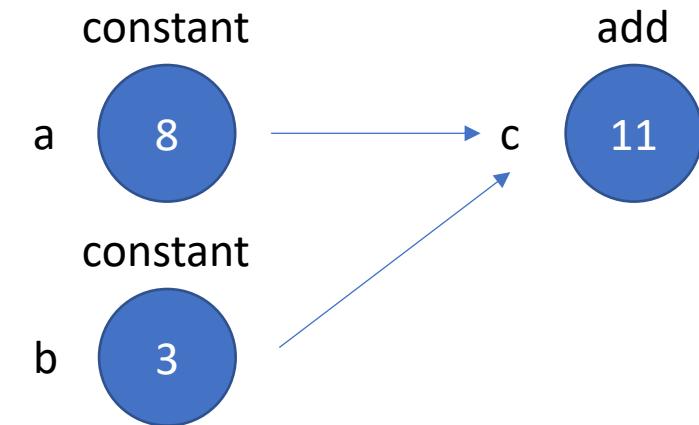






TensorFlow

```
import tensorflow as tf  
a = tf.constant(8)  
b = tf.constant(3)  
c = tf.add(a, b) # or a + b, operator is overloaded
```



```
Last login: Tue Apr 10 09:40:34 on ttys001
douglasmbpr:~ douglasstarnes$ ipython
Python 3.6.3 |Anaconda custom (64-bit)| (default, Oct  6 2017, 12:04:38)
Type 'copyright', 'credits' or 'license' for more information
IPython 6.1.0 -- An enhanced Interactive Python. Type '?' for help.
```

```
In [1]: import tensorflow as tf
```

```
In [2]: a = tf.constant(8)
```

```
In [3]: b = tf.constant(3)
```

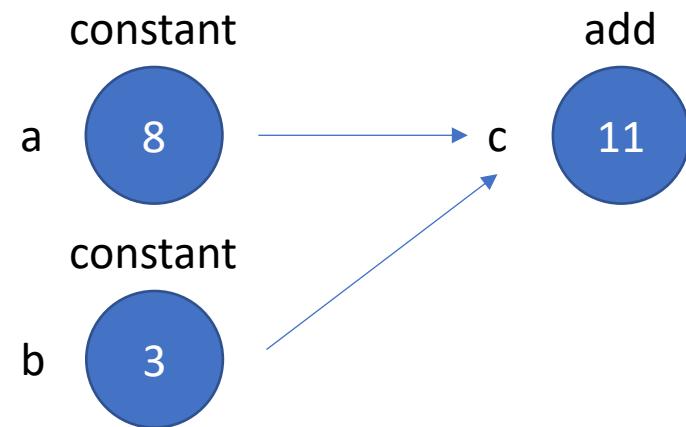
```
In [4]: c = tf.add(a, b)
```

```
In [5]: print(c)
```

```
Tensor("Add:0", shape=(), dtype=int32)
```

```
In [6]: 
```

```
import tensorflow as tf  
a = tf.constant(8)  
b = tf.constant(3)  
c = tf.add(a, b) # or a + b  
with tf.Session() as s:  
    print(s.run(c))
```



```
In [5]: print(c)
Tensor("Add:0", shape=(), dtype=int32)
```

```
In [6]: with tf.Session() as s:
...:     print(s.run(c))
...:
```

W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow library wasn't compiled to use SSE4.1 instructions, but these are available on your machine and could speed up CPU computations.

W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow library wasn't compiled to use SSE4.2 instructions, but these are available on your machine and could speed up CPU computations.

W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow library wasn't compiled to use AVX instructions, but these are available on your machine and could speed up CPU computations.

W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow library wasn't compiled to use AVX2 instructions, but these are available on your machine and could speed up CPU computations.

W tensorflow/core/platform/cpu_feature_guard.cc:45] The TensorFlow library wasn't compiled to use FMA instructions, but these

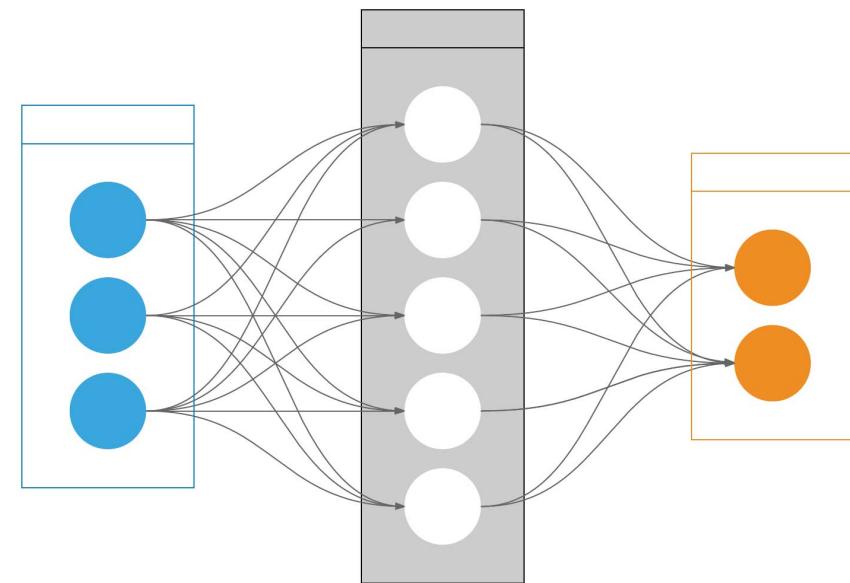
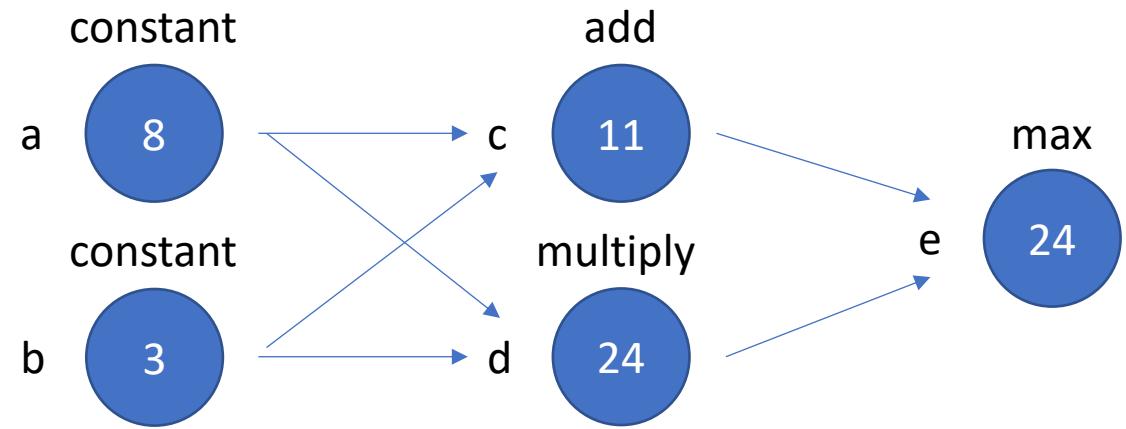
11

```
In [7]:
```

```

import tensorflow as tf
a = tf.constant(8)
b = tf.constant(3)
c = tf.add(a, b) # or a + b
with tf.Session() as s:
    print(s.run(c))
d = tf.multiply(a, b) # or a * b
e = tf.max(c, d)
with tf.Session() as s:
    print(s.run(e))

```



Meet MNIST

Scans of handwritten digits

Each image is 28x28 pixels

Each pixel value is 0 (black)
to 255 (white)

Feature vector is the entire
image (784 pixels)

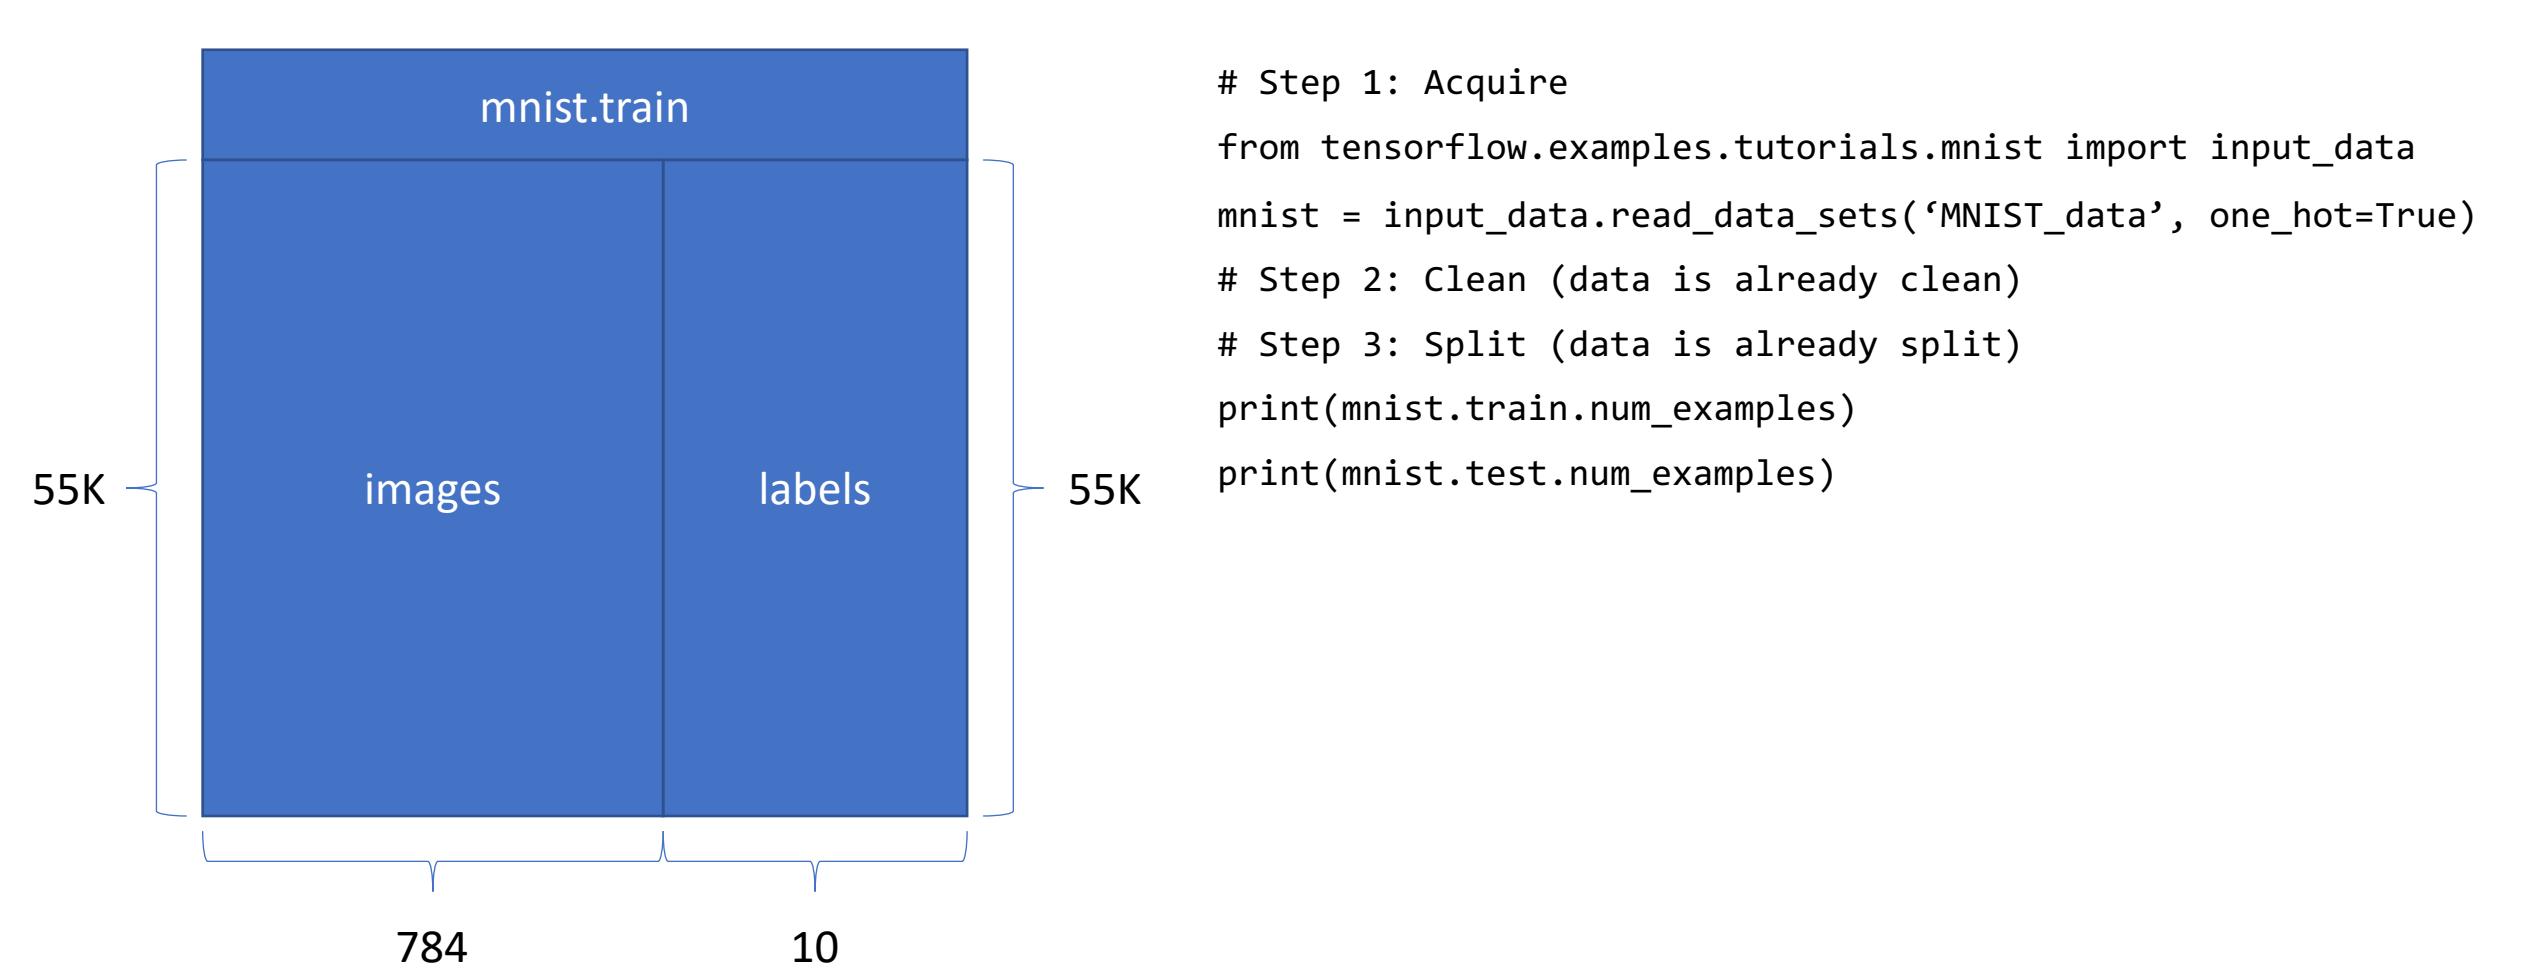
The target is the digit
represented (0-9)

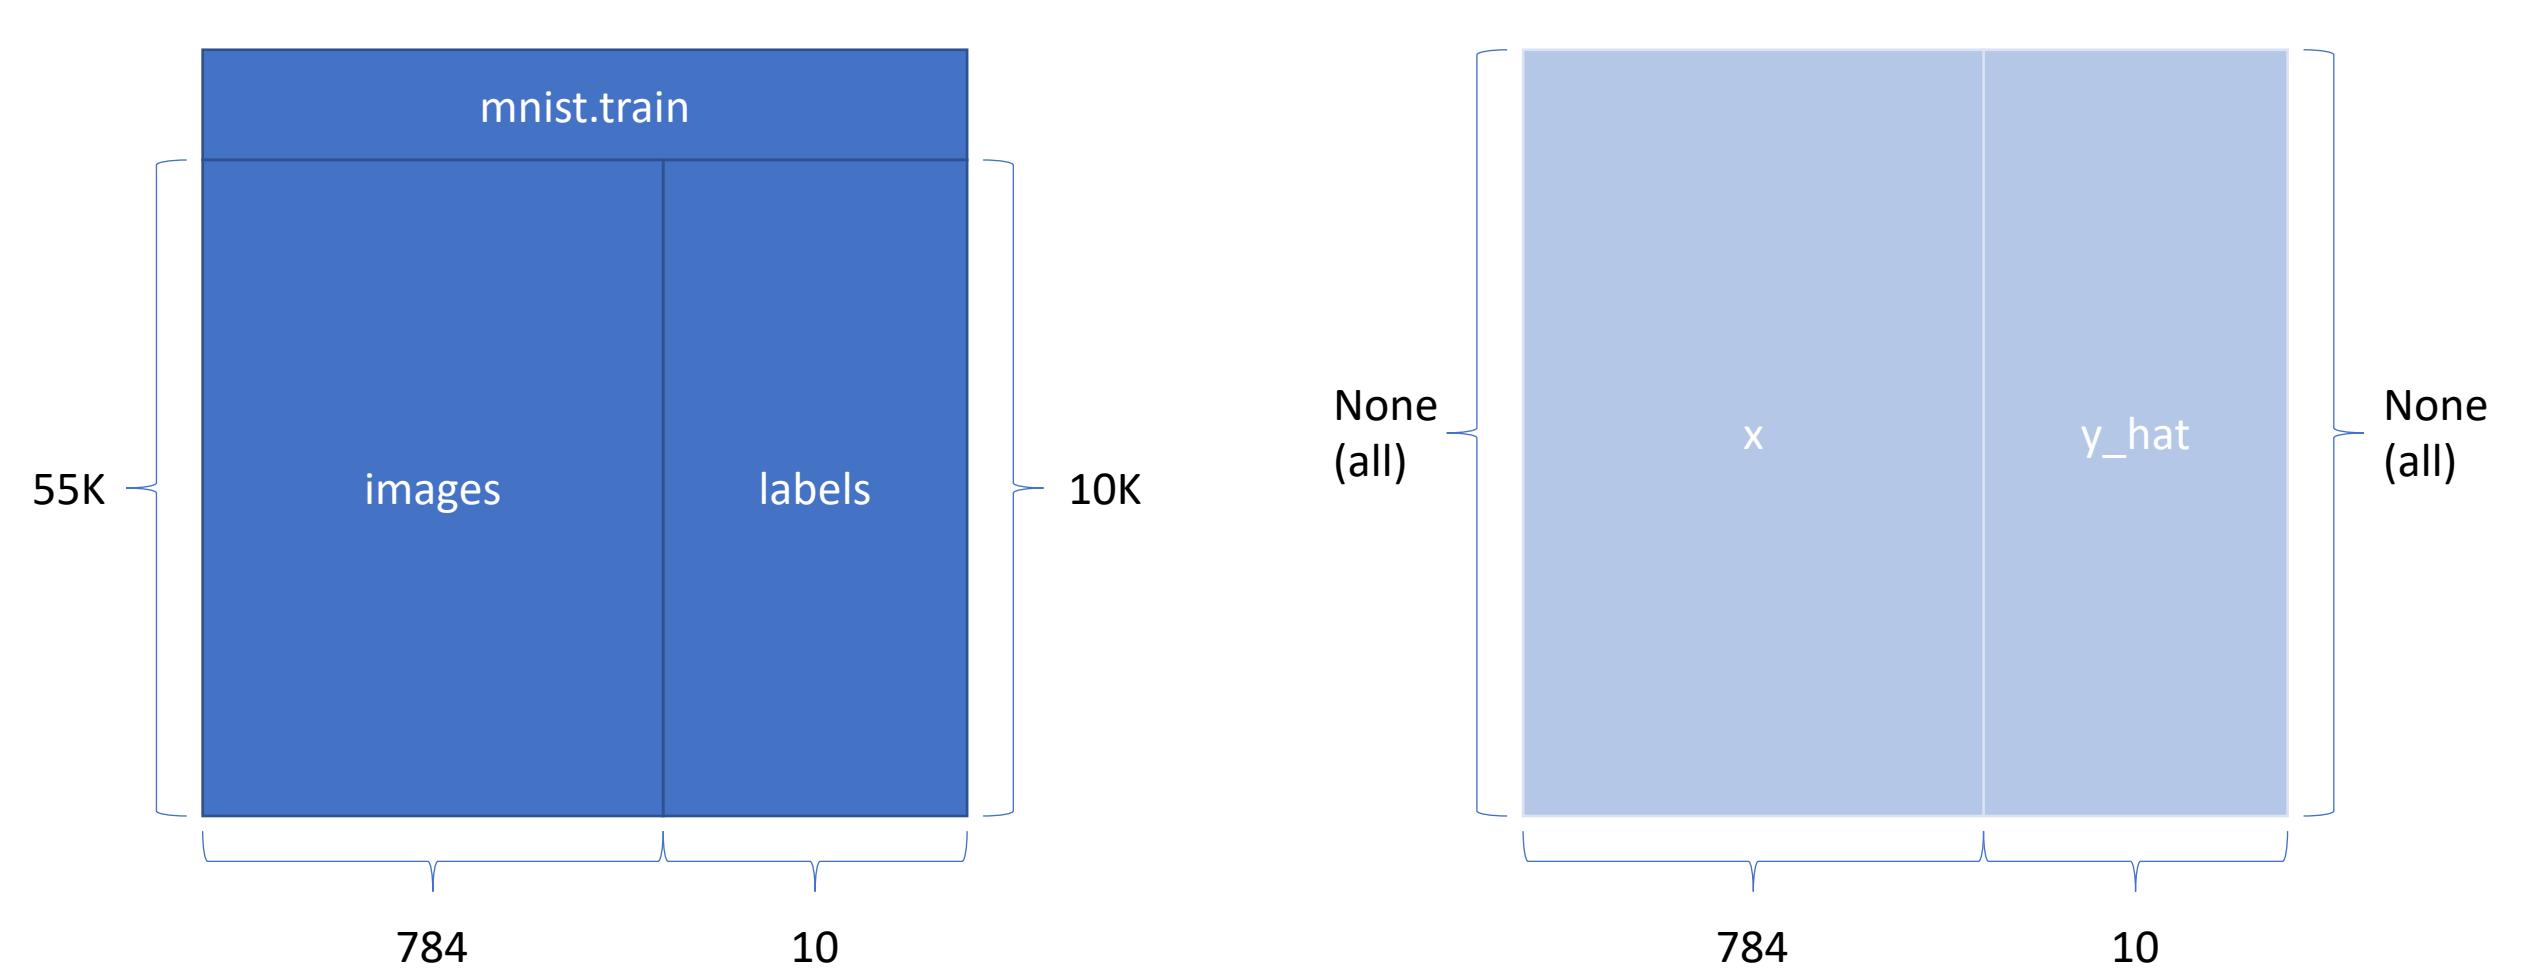
Interesting variations
(ie. fashion MNIST)



```
# Step 1: Acquire
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('MNIST_data', one_hot=True)
# Step 2: Clean (data is already clean)
# Step 3: Split (data is already split)
print(mnist.train.num_examples) # 55000
print(mnist.test.num_examples) # 10000
# Step 4: Train
x = tf.placeholder(tf.float32, shape=[None, 784]) # input for feature vectors
y_hat = tf.placeholder(tf.float32, shape=[None, 10]) # input for targets
weights = tf.Variable(tf.truncated_normal([784,10])) # initialize model with random values
bias = tf.Variable(tf.truncated_normal([10])) # initialize model with random values
classifier = tf.matmul(x, w) + b # prediction
y = tf.nn.softmax(classifier) # activation function
cost_function =
    tf.reduce_mean(-tf.reduce_sum(y_hat * tf.log(y))) # measure of similiarity
```

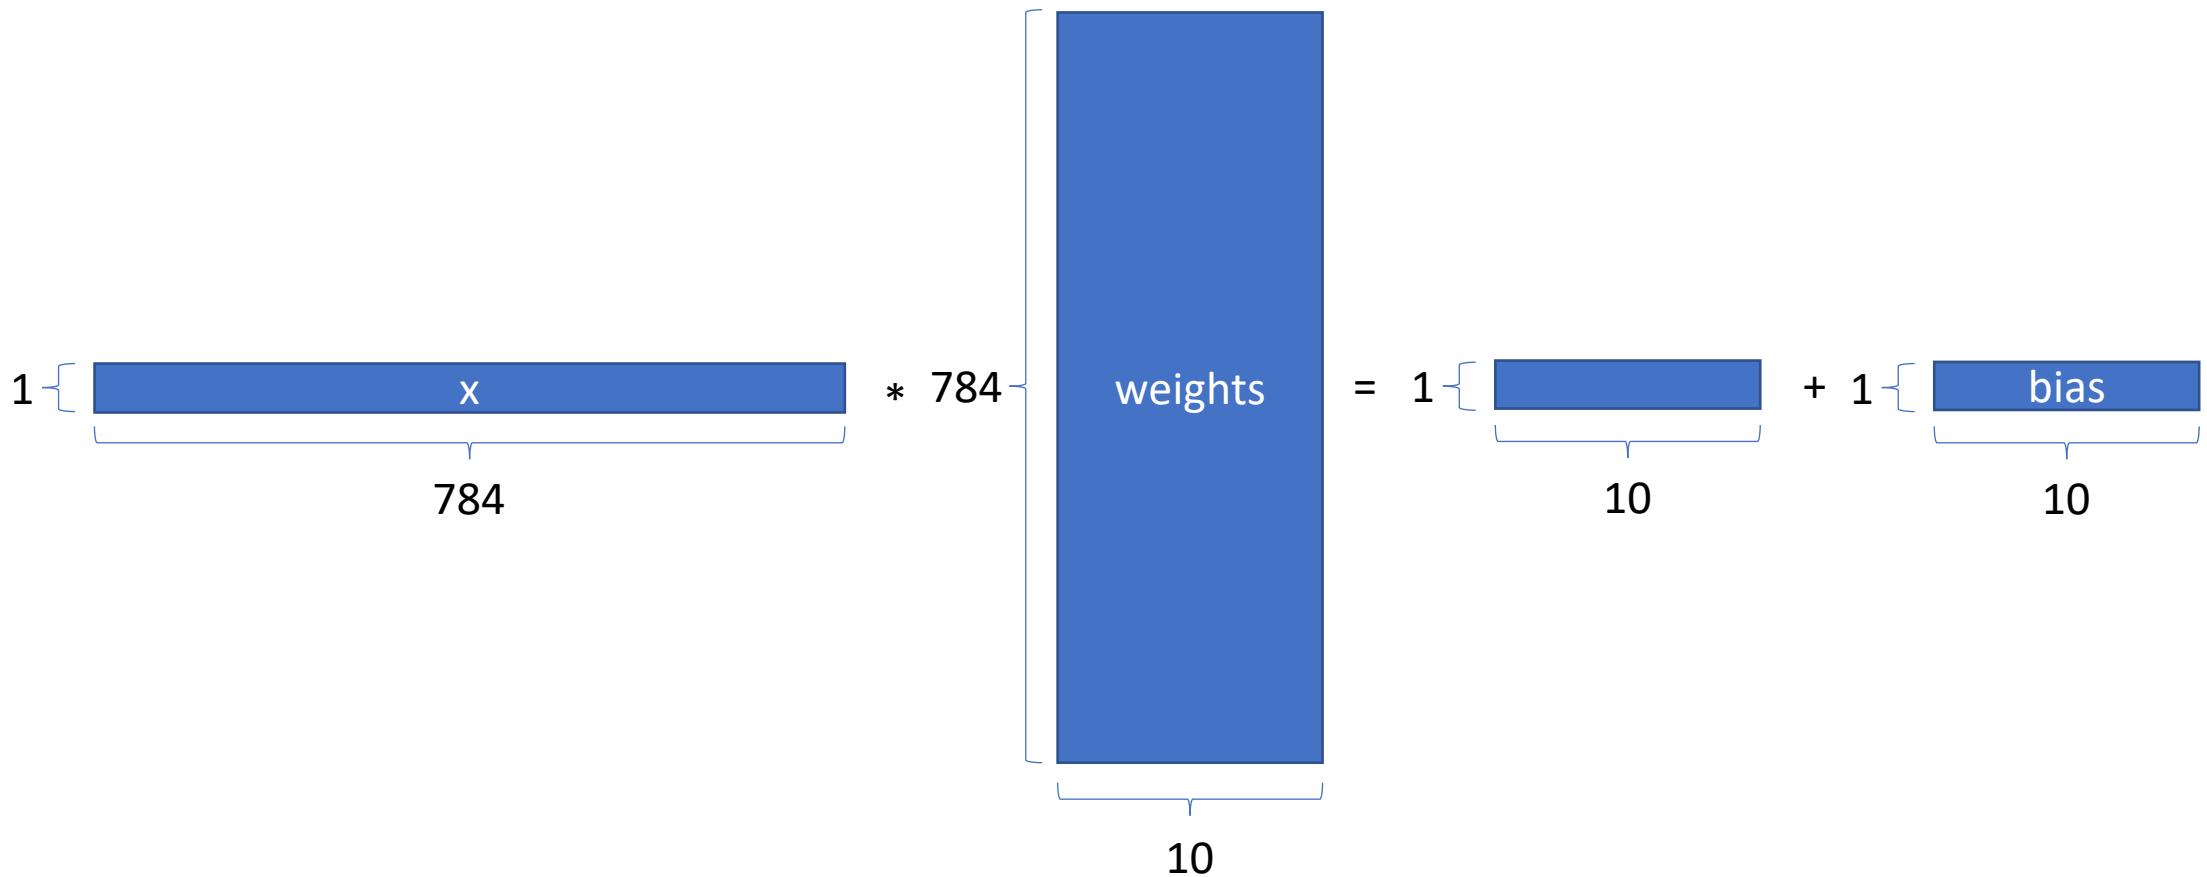
```
optimizer = tf.train.GradientDescentOptimizer(0.5).minimize(cost_function) # optimizer  
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer()) # assign random values  
    for _ in mnist.train.num_examples / 100: # run one epoch  
        batch = mnist.train.next_batch(100)  
        sess.run(optimizer, feed_dict={x: batch[0], y_hat: batch[1]})  
  
# Step 5: Test  
correct_predictions = tf.equal(tf.argmax(y, 1), tf.argmax(y_hat, 1)) # total no. correct  
  
# Step 6: Score  
average_correct = tf.reduce_mean(tf.cast(correct_predictions, tf.float32)) # avg. no. correct  
with tf.Session() as sess:  
    print(  
        sess.run(average_correct,  
            feed_dict={x: mnist.test.images, y_hat: mnist.test.labels})  
  
# Step 7: Evaluate - make some pretty pictures
```

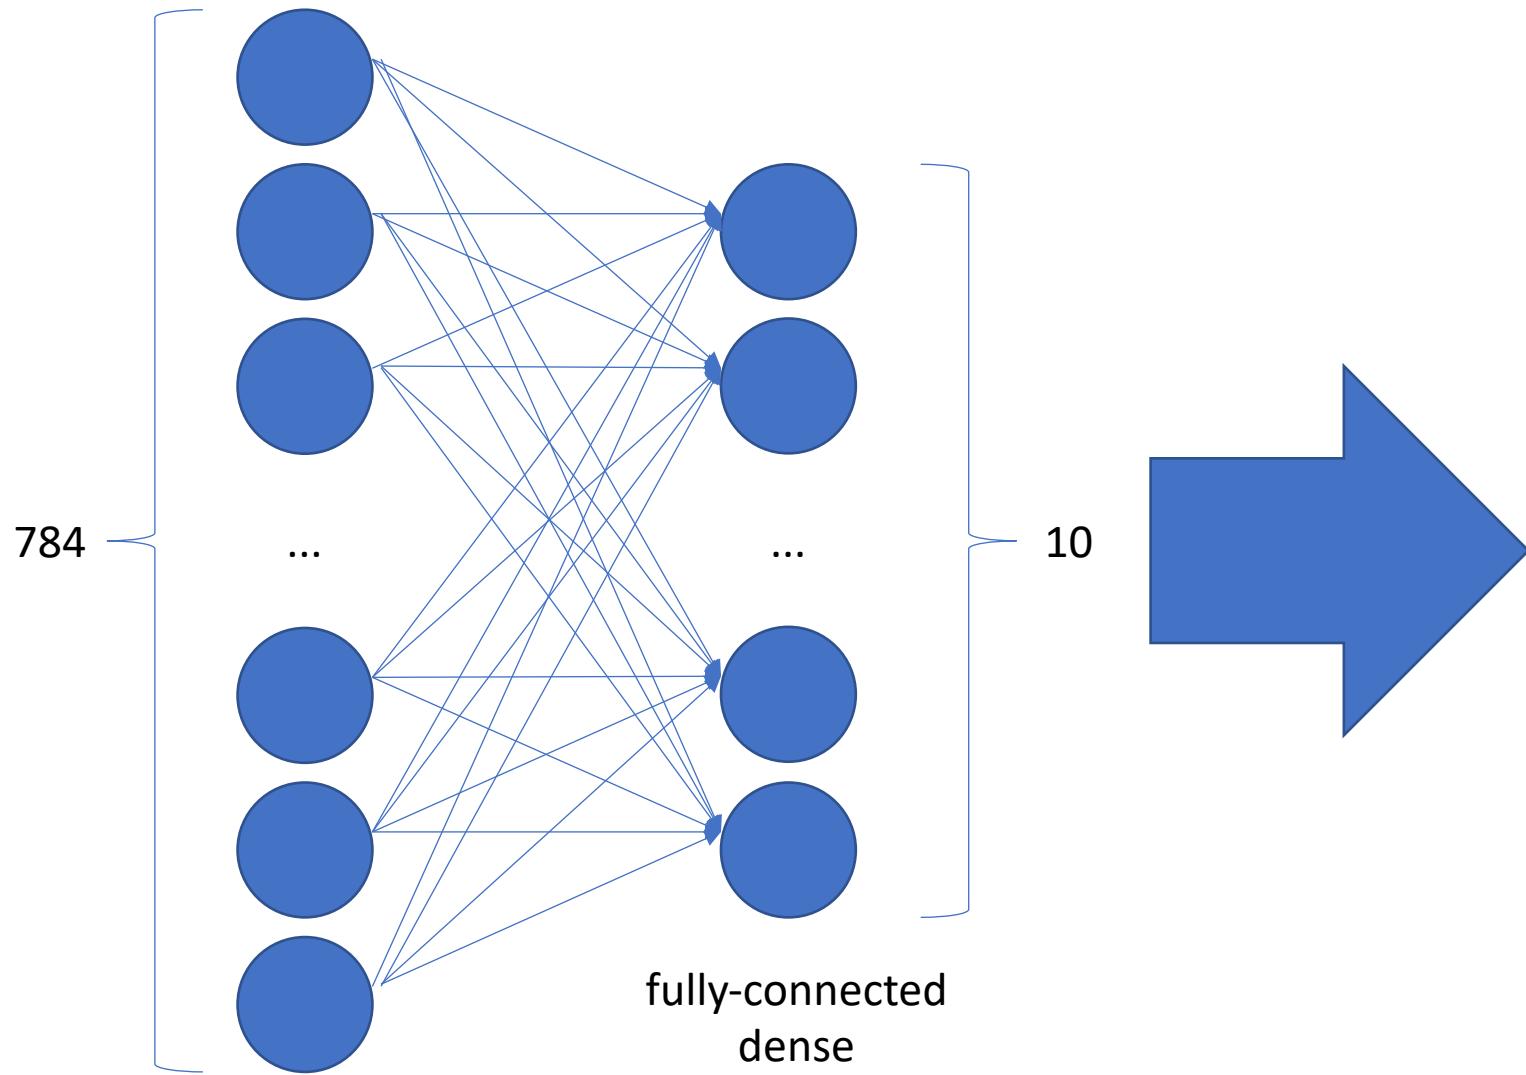




```
# Step 4: Train  
x = tf.placeholder(tf.float32, shape=[None, 784])  
y_hat = tf.placeholder(tf.float32, shape=[None, 10])
```

```
weights = tf.Variable(tf.truncated_normal([784,10])  
bias = tf.Variable(tf.truncated_normal([10]))  
classifier = tf.matmul(x, weights) + bias
```





```
y = tf.nn.softmax(classifier)
```

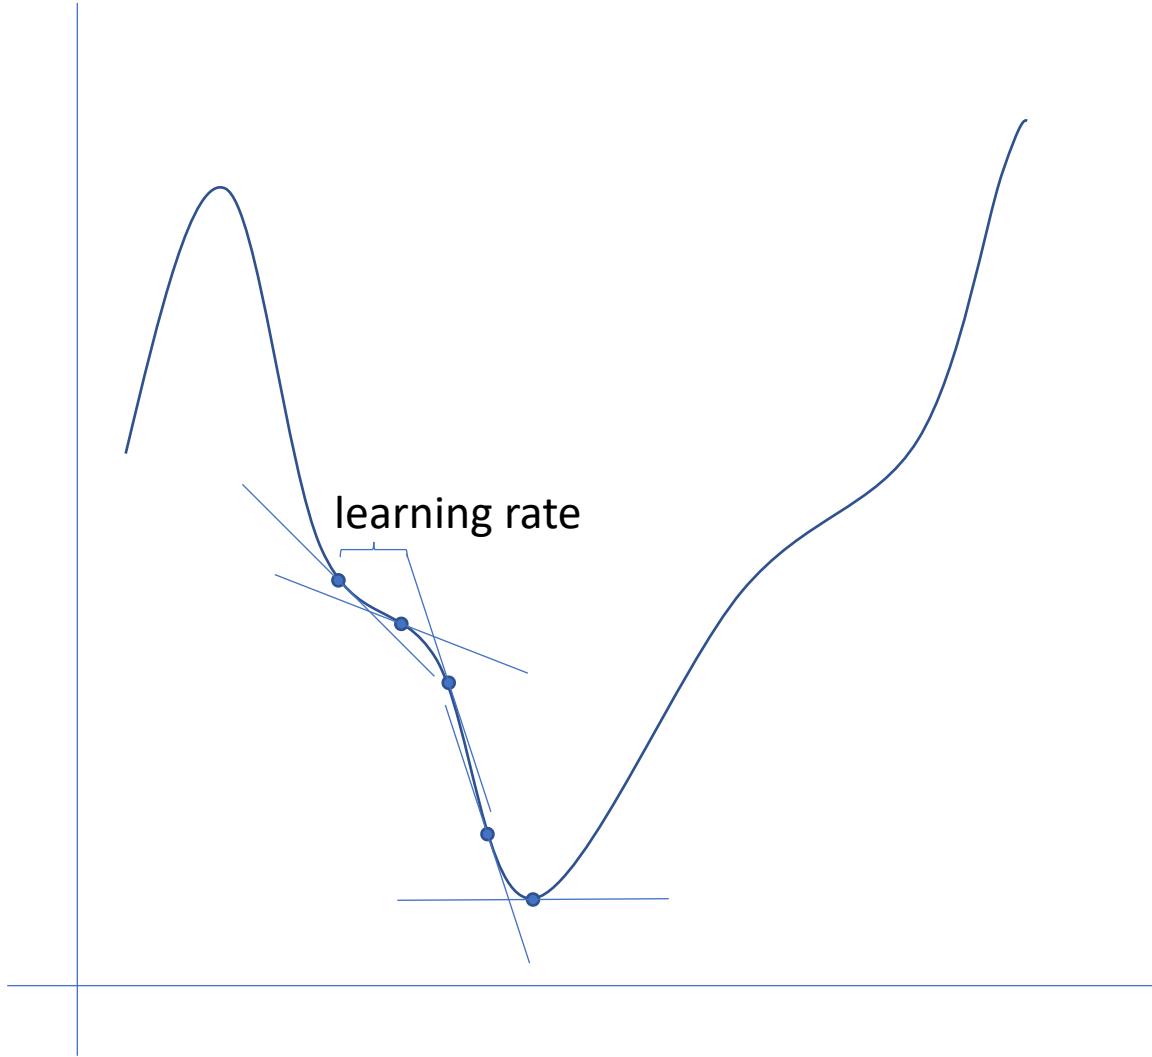
$$\text{softmax}(i_x) = \frac{e^{i_x}}{\sum_{j=1}^k e^{i_j}}$$

```
from math import exp
i = [100, 101, 102]
def softmax(x):
    return exp(i[x]) / sum([exp(j) for j in i])
print([softmax(a) for a in range(len(i))])
```

```
[0.09003057317038046, 0.24472847105479764, 0.6652409557748219]
```

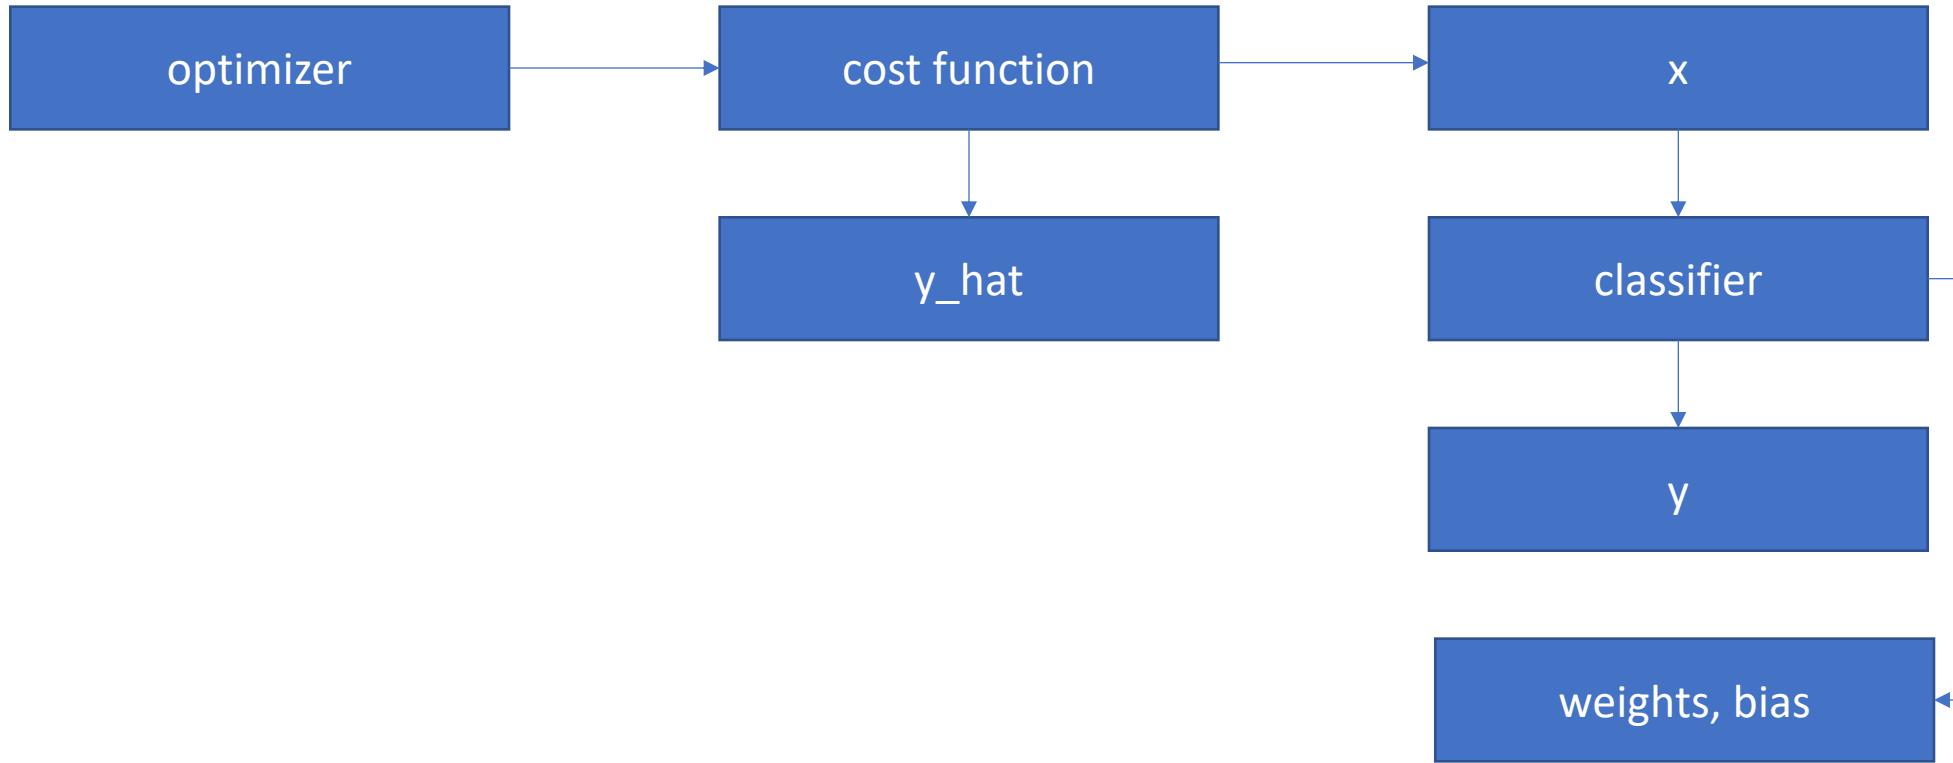
9% 24% 67%

```
optimizer = tf.train.GradientDescentOptimizer(0.5).minimize(cost_function)
```



If the learning rate is too large, we will skip over the minimum.
If the learning rate is too small, we will waste time.

```
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer())  
    for _ in mnist.train.num_examples / 100:  
        batch = mnist.train.next_batch(100)  
        sess.run(optimizer, feed_dict={x: batch[0], y_hat: batch[1]})
```



```
correct_predictions = tf.equal(tf.argmax(y, 1), tf.argmax(y_hat, 1))
```

0	1	2	3	4	5	6	7	8	9
.003	.24	.9	.012	.004	0	.031	.001	.013	.008
.01	.019	.012	.003	.9	.02	.012	.014	.005	0

y

0	1	2	3	4	5	6	7	8	9
0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	0	0

y_hat

```
average_correct = tf.reduce_mean(tf.cast(correct_predictions, tf.float32))
```

[True, True, False, True, True, True, False, True, True, True]

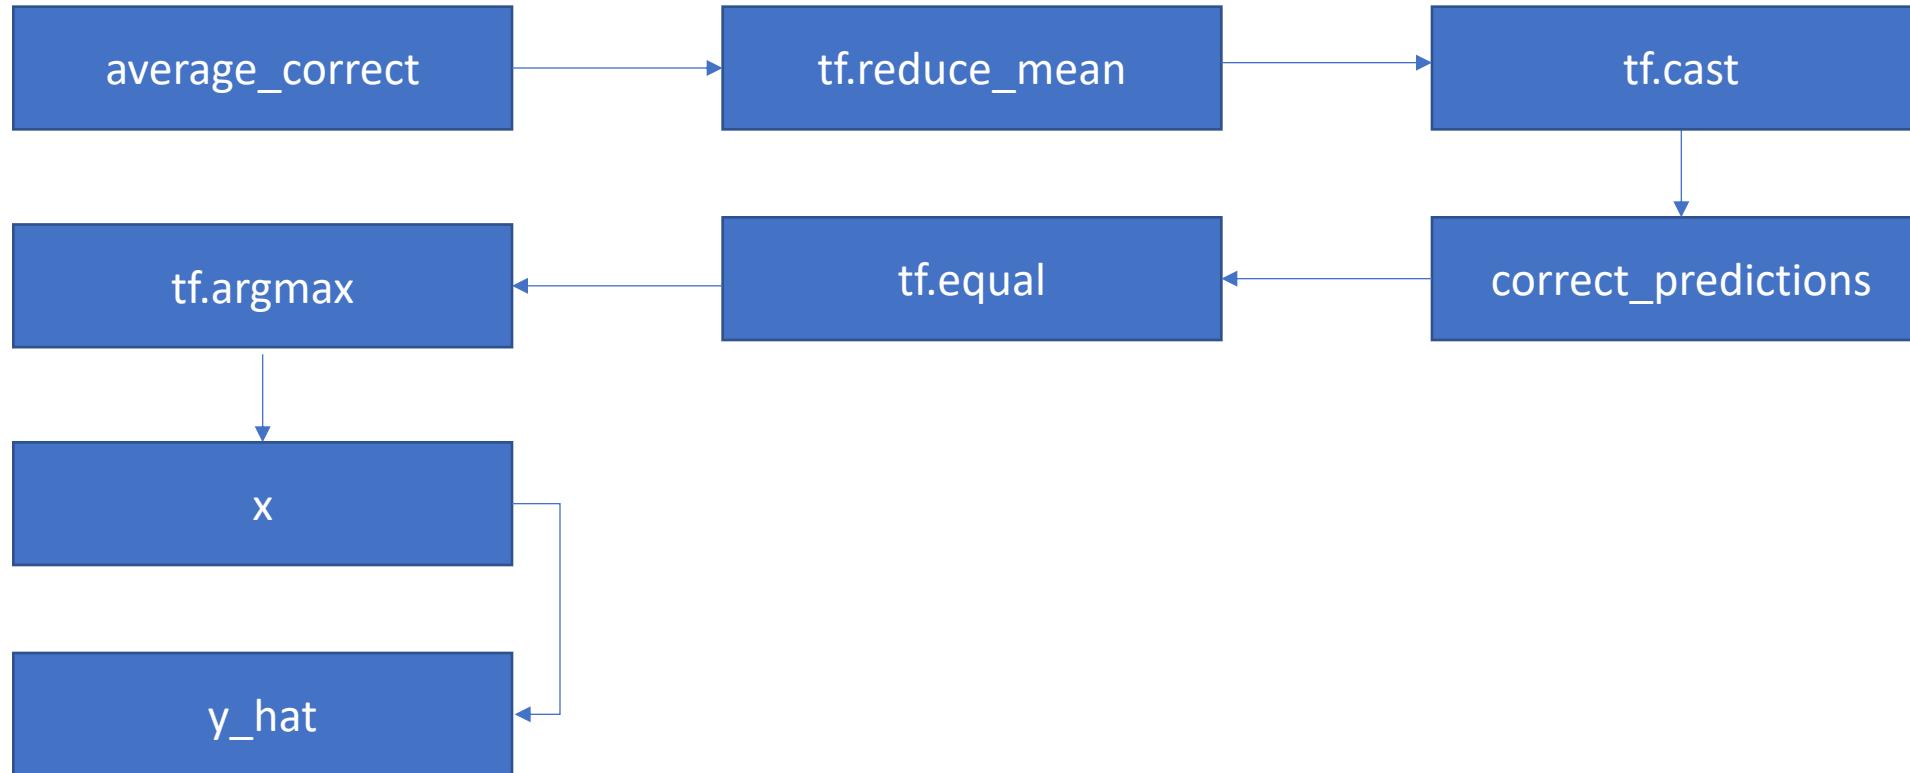
tf.cast

[1.0, 1.0, 0.0, 1.0, 1.0, 1.0, 0.0, 1.0, 1.0, 1.0]

tf.reduce_mean

0.8

```
with tf.Session() as sess:  
    sess.run(average_correct, feed_dict={x: mnist.test.images,  
                                         y_hat: mnist.test.labels})
```



DEMO





That works but surely there
is an easier method?



Keras

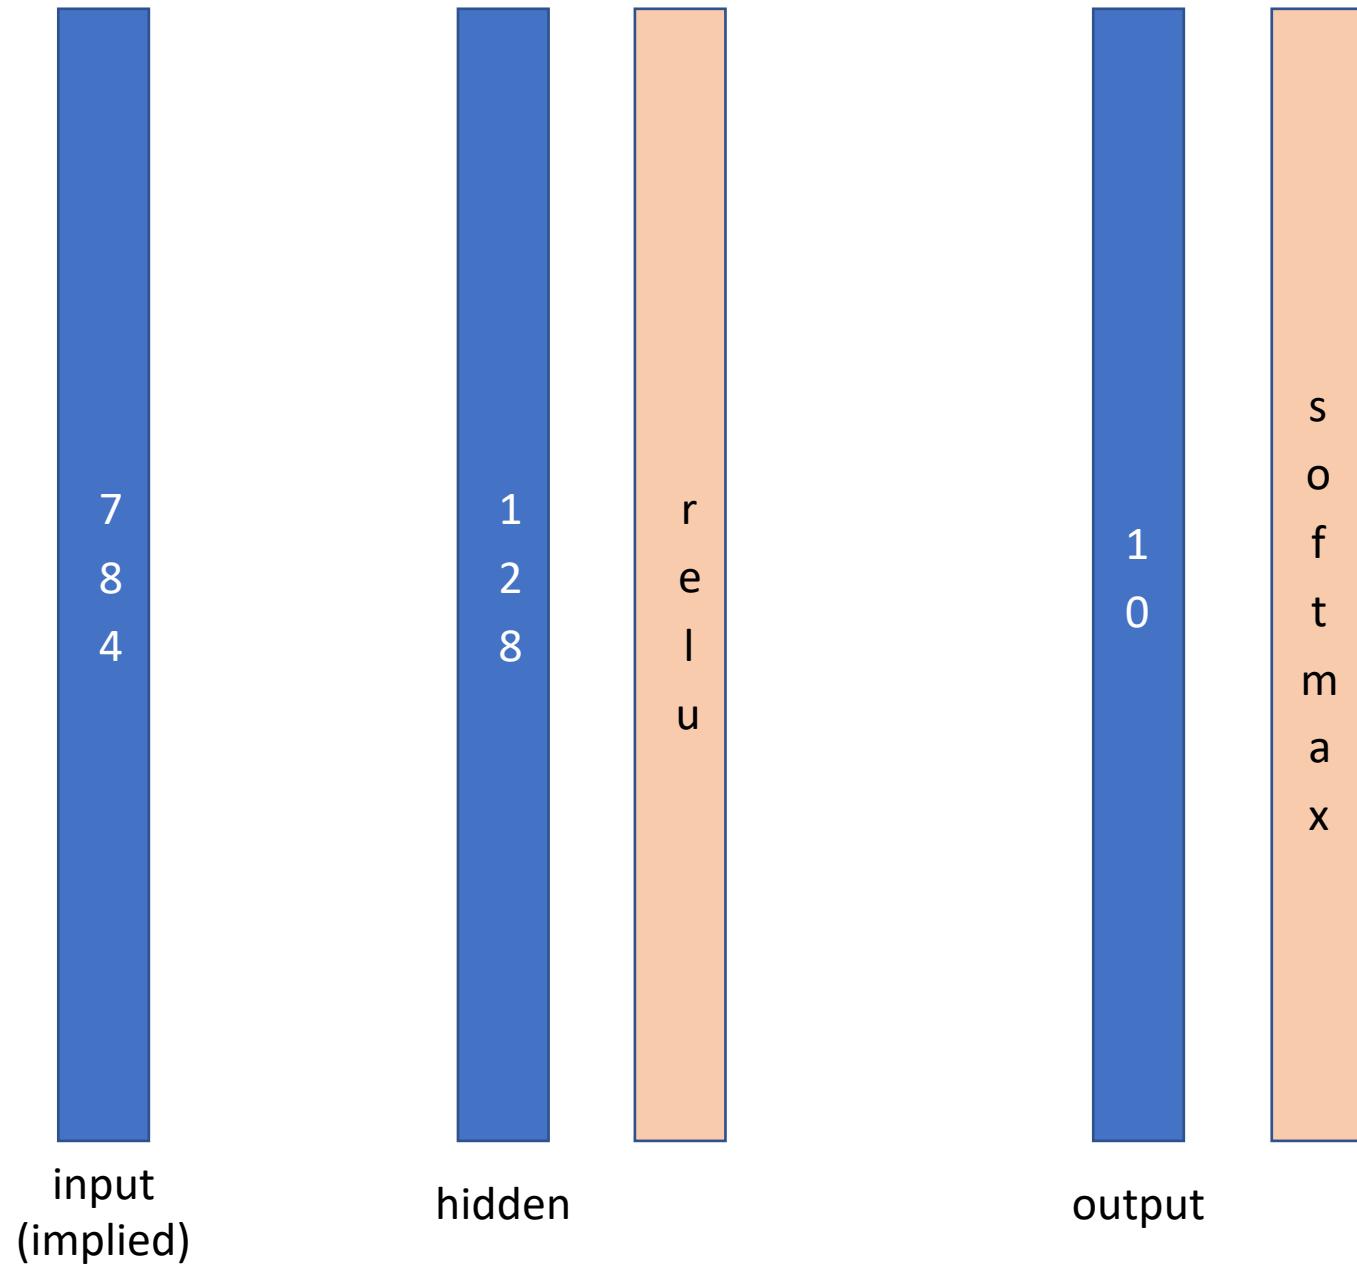


```
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers.core import Dense, Activation
from keras.optimizers import SGD
from keras.utils import np_utils

(X_train, y_train), (X_test, y_test) = mnist.load_data()
X_train = X_train.reshape(60000, 784).astype('float32') / 255 # apply to X_test
Y_train = np_utils.to_categorical(y_train, 10) # apply to y_test
model = Sequential()

model.add(Dense(128, input_shape=(784,)))
model.add(Activation('relu'))
model.add(Dense(10))
model.add(Activation('softmax'))

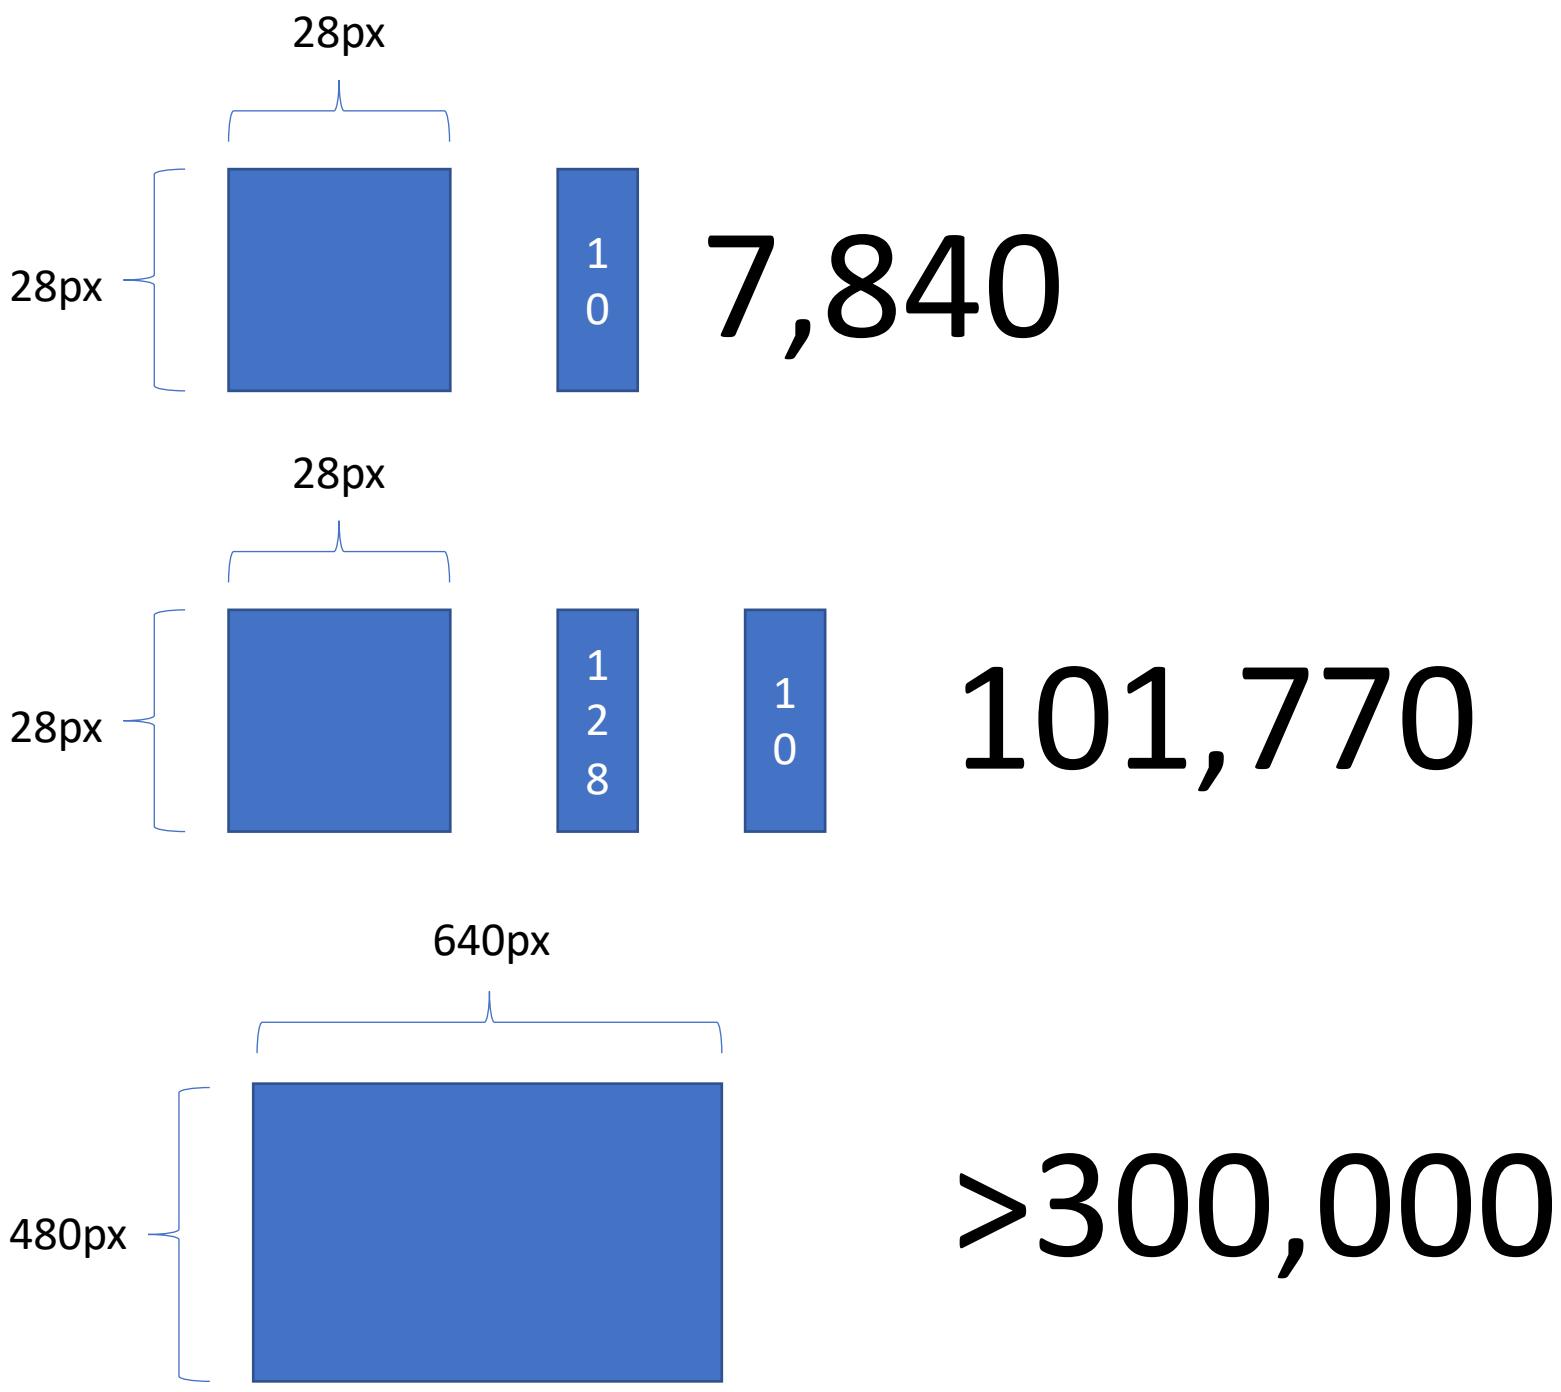
model.compile(loss='categorical_crossentropy', optimizer=SGD(), metrics=[ 'accuracy' ])
model.fit(X_train, Y_train, batch_size=128, epochs=20) # around 94%
score = model.evaluate(X_test, Y_test)
```



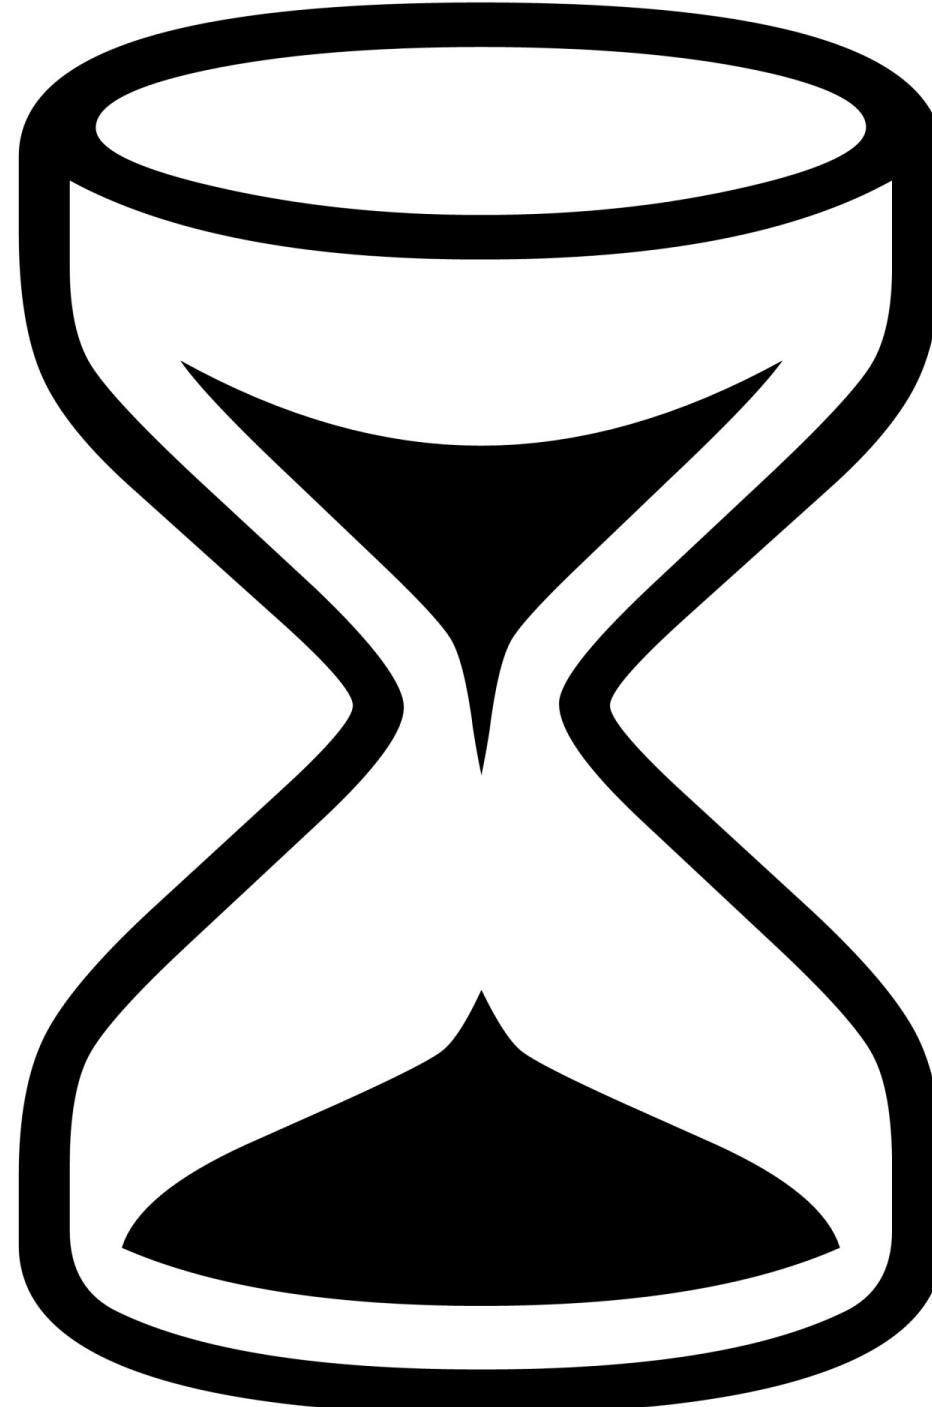


That's all there is to it?

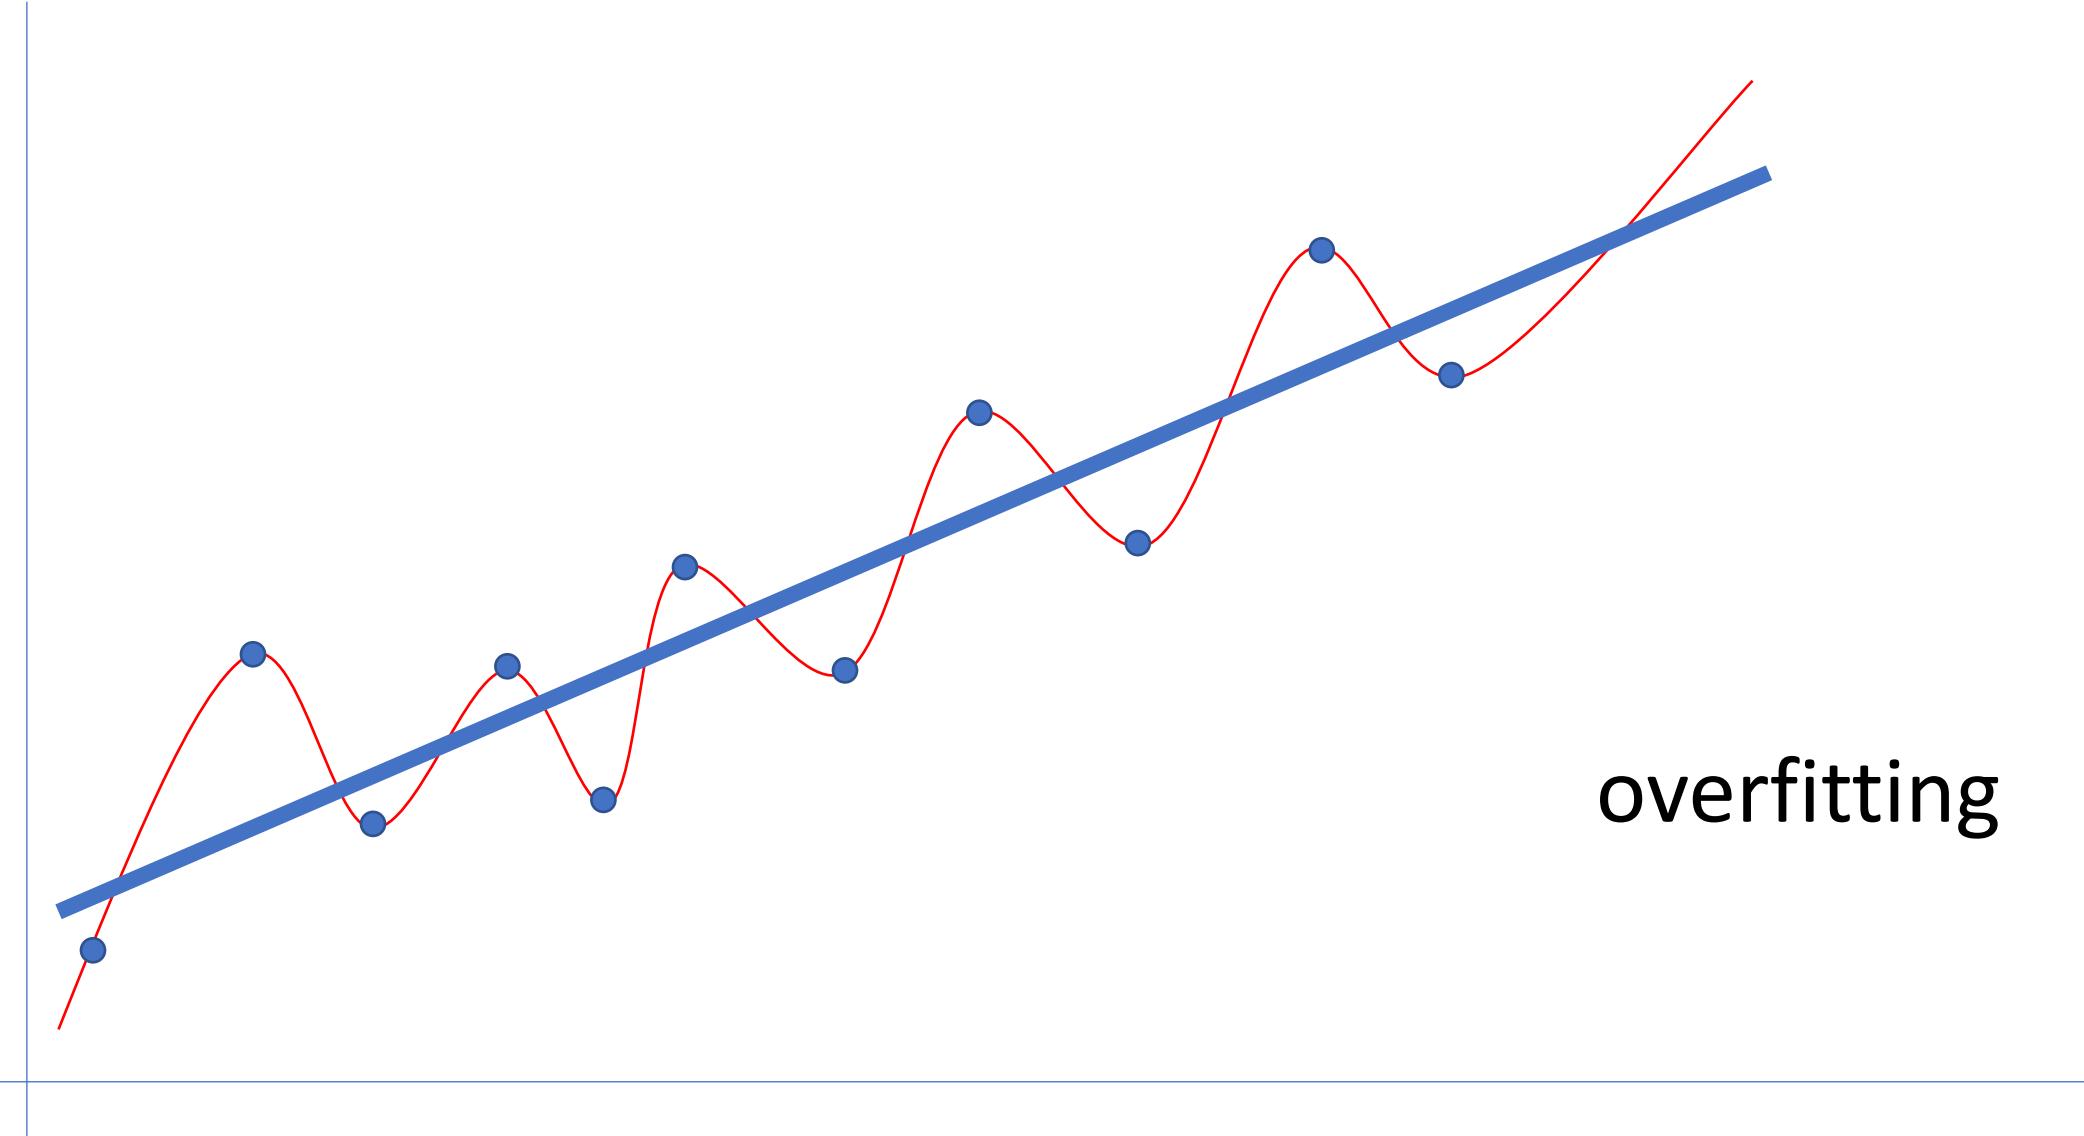
Surely there is a catch.



@poweredbyaltnet



@poweredbyaltnet





@poweredbyaltnet

Goal: consider fewer, more relevant, features

How to determine which features are relevant?

Let the network decide!

Convolutional Neural Network (CNN)

Three new layers

Convolutional layers – detect features

Pooling layers – summarize features

Dropout – prevent overfitting

Features are then used in dense layers

i
n
p
u
t

c
o
n
v
2
D

1	0	0	1	1	0
1	1	1	0	1	1
0	1	1	1	0	1
1	1	1	0	1	1
0	0	1	1	0	1
0	1	1	0	0	0

1	0	0
1	1	1
0	1	1

filter/kernel

1 1	0 0	0 0	1	1	0
1 1	1 1	1 1	0	1	1
0 0	1 1	1 1	1	0	1
1	1	1	0	1	1
0	0	1	1	0	1
0	1	1	0	0	0

6

1	0 1	0 0	1 0	1	0
1	1 1	1 1	0 1	1	1
0	1 0	1 1	1 1	0	1
1	1	1	0	1	1
0	0	1	1	0	1
0	1	1	0	0	0

6	4
---	---

1	0	0 1	1 0	1 0	0
1	1	1 1	0 1	1 1	1
0	1	1 0	1 1	0 1	1
1	1	1	0	1	1
0	0	1	1	0	1
0	1	1	0	0	0

6	4	3
---	---	---

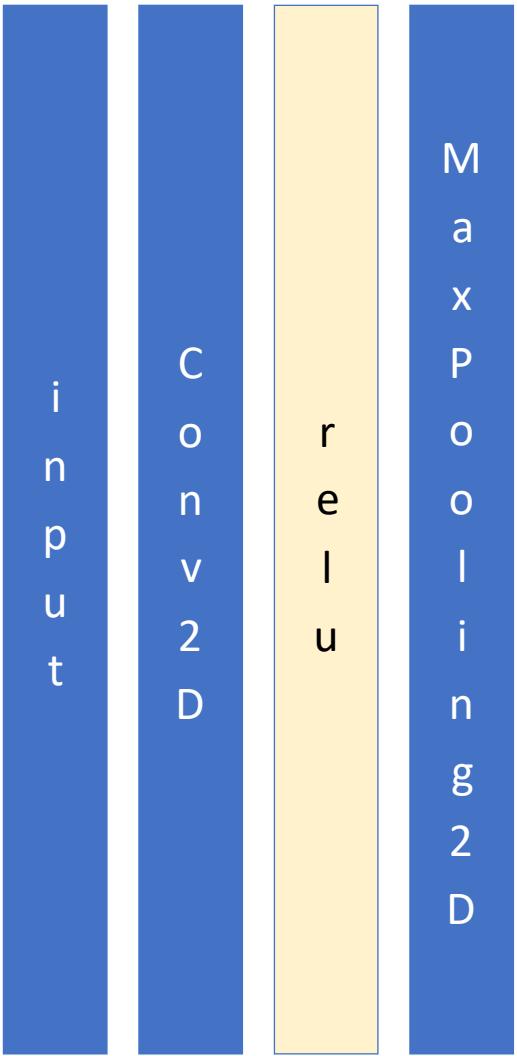
1	0	0	1 1	1 0	0 0
1	1	1	0 1	1 1	1 1
0	1	1	1 0	0 1	1 1
1	1	1	0	1	1
0	0	1	1	0	1
0	1	1	0	0	0

6	4	3	4
---	---	---	---

1	0	0	1	1	0
1	1	1	0	1	1
0	1	1	1	0	1
1	1	1	0 1	1 0	1 0
0	0	1	1 1	0 1	1 1
0	1	1	0 0	0 1	0 1

6	4	3	4
5	5	4	4
4	5	4	4
4	4	3	2

convolution



6	4	3	4
5	5	4	4
4	5	4	4
4	4	3	2

6

6	4	3	4
5	5	4	4
4	5	4	4
4	4	3	2

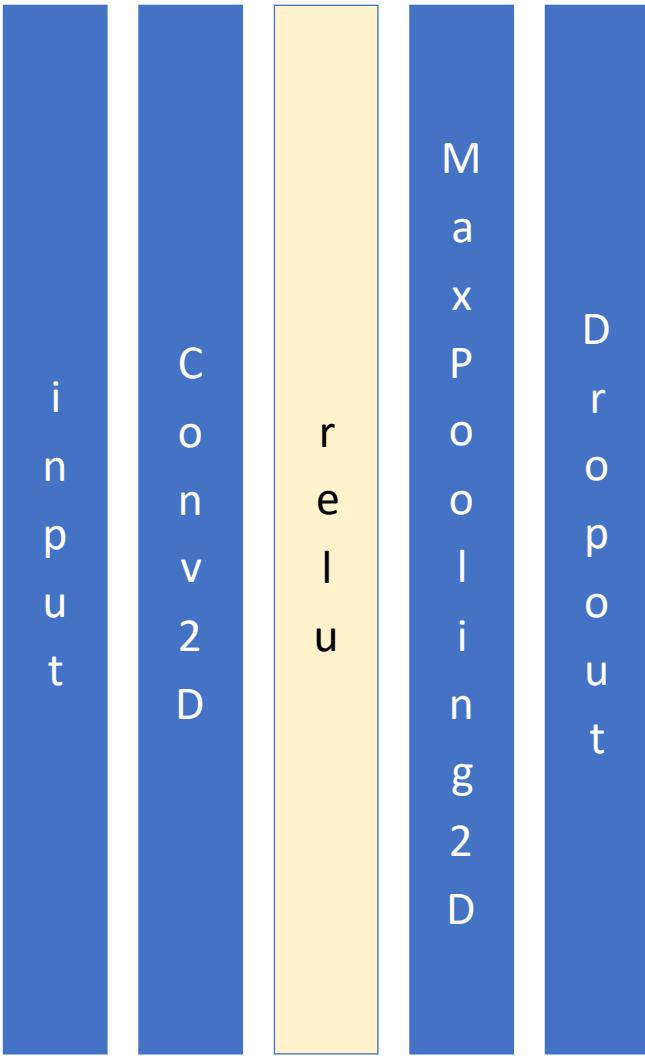
6	4
---	---

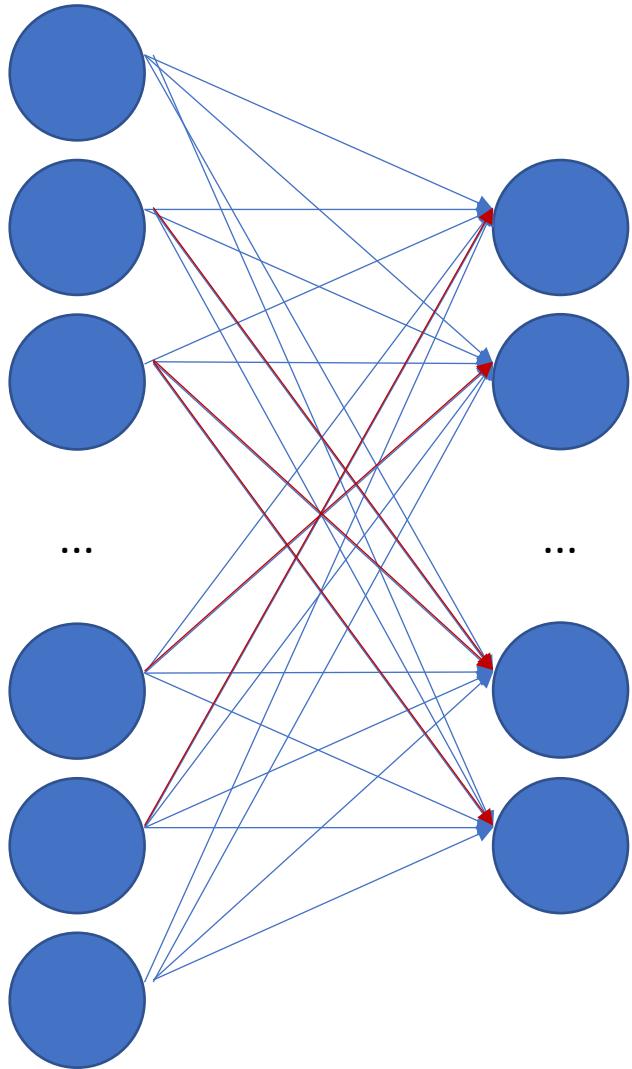
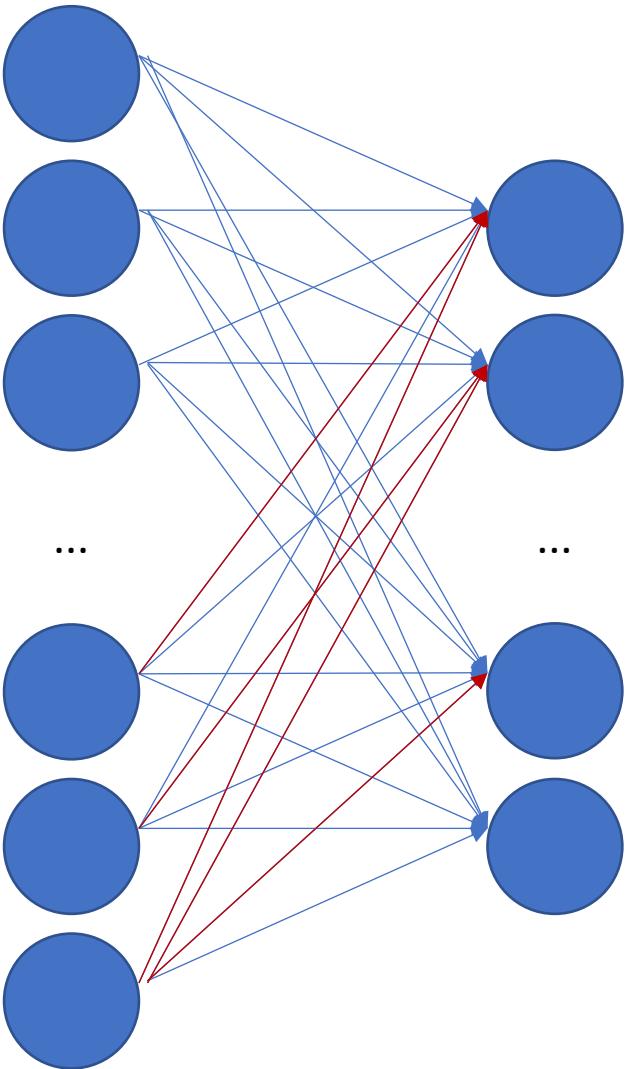
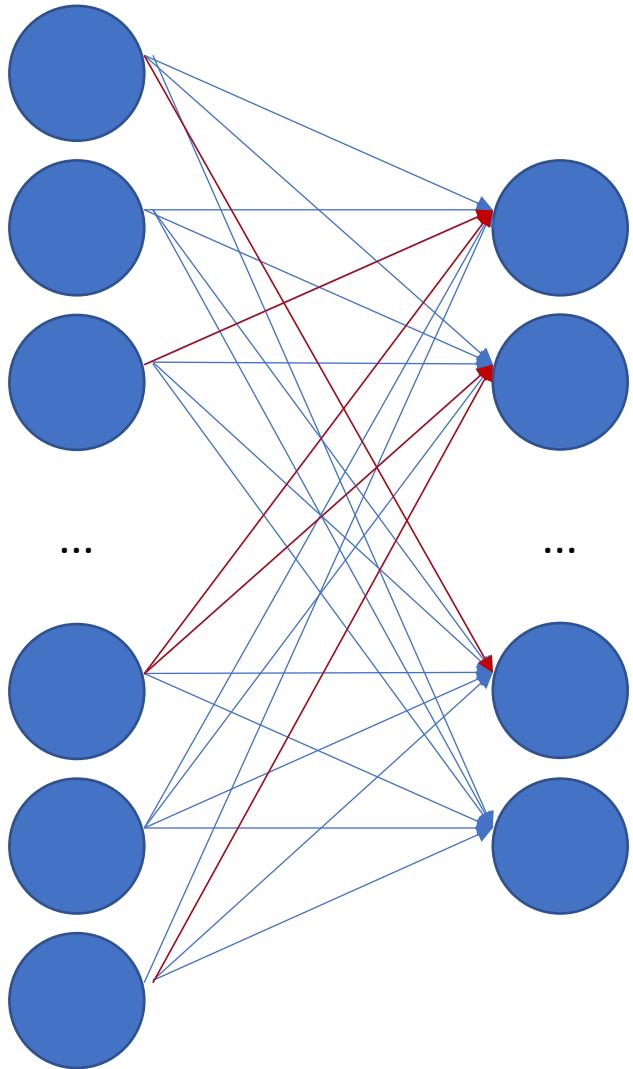
6	4	3	4
5	5	4	4
4	5	4	4
4	4	3	2

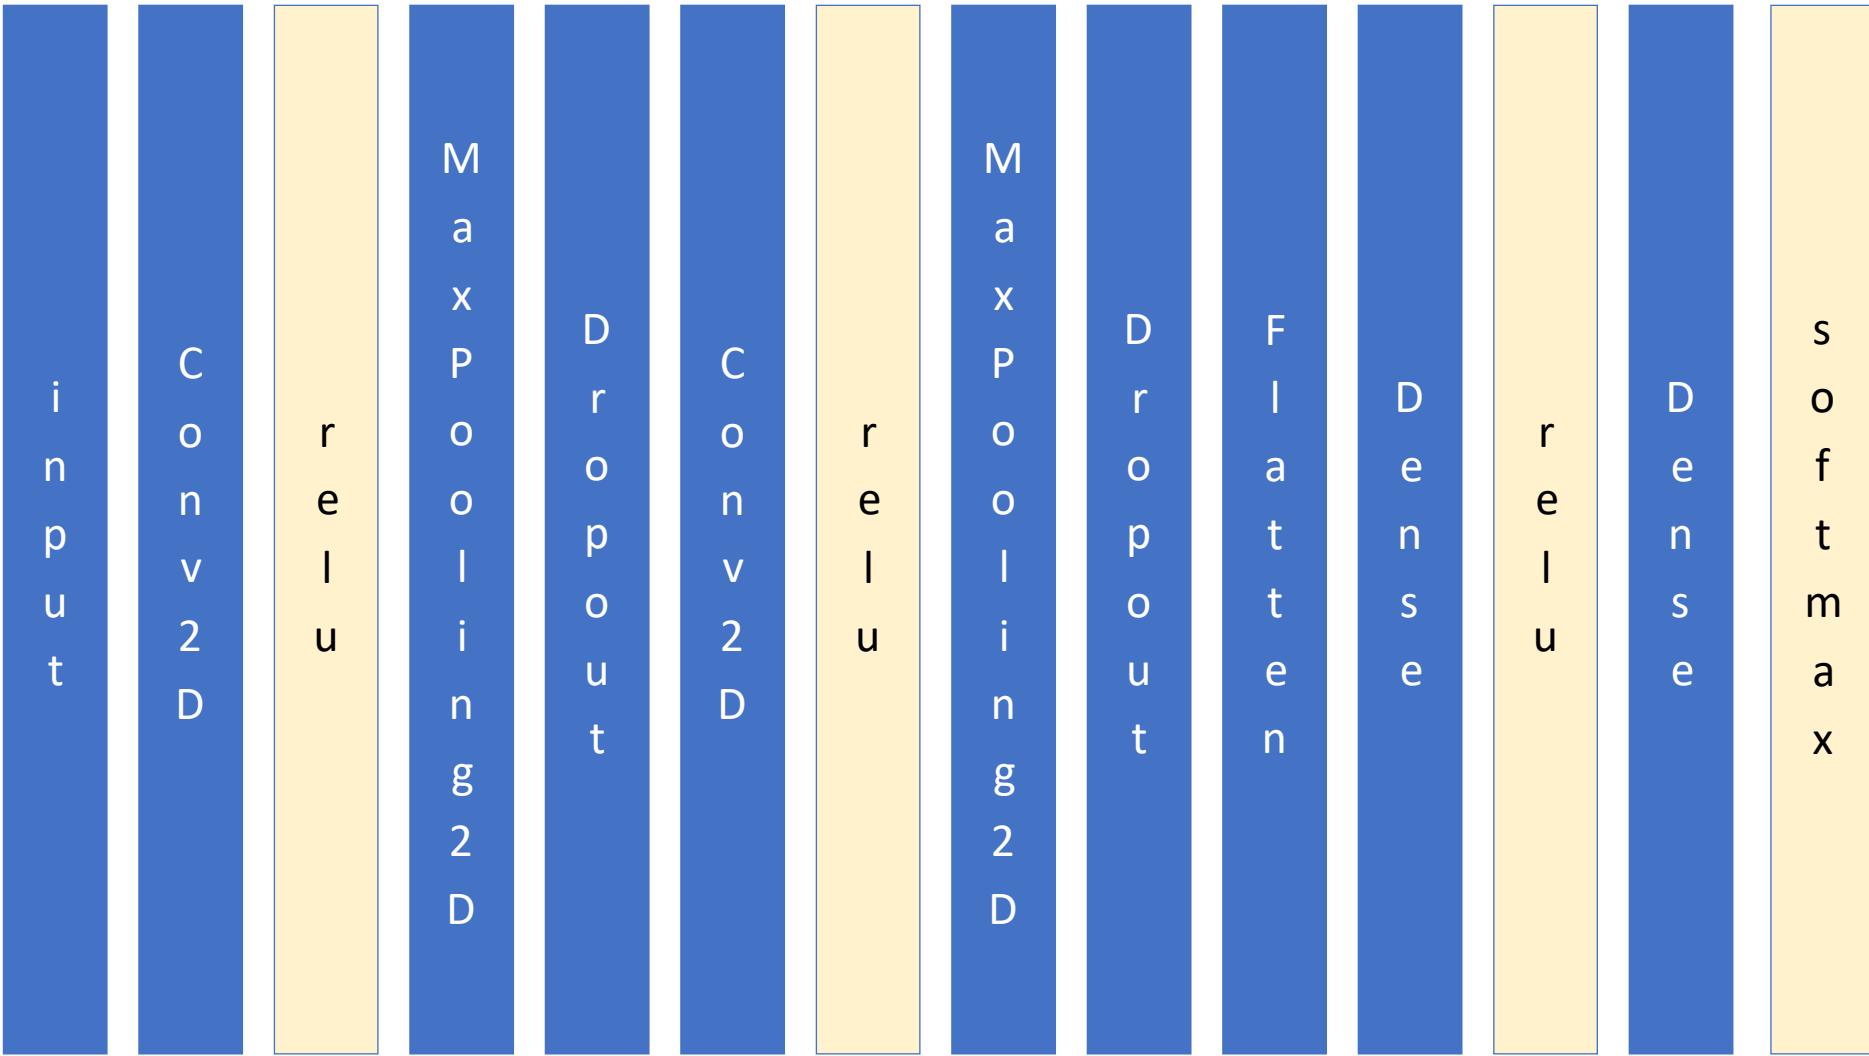
6	4
5	

6	4	3	4
5	5	4	4
4	5	4	4
4	4	3	2

6	4
5	4







```
from keras.datasets import cifar10

from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten,
Conv2D, MaxPooling2D

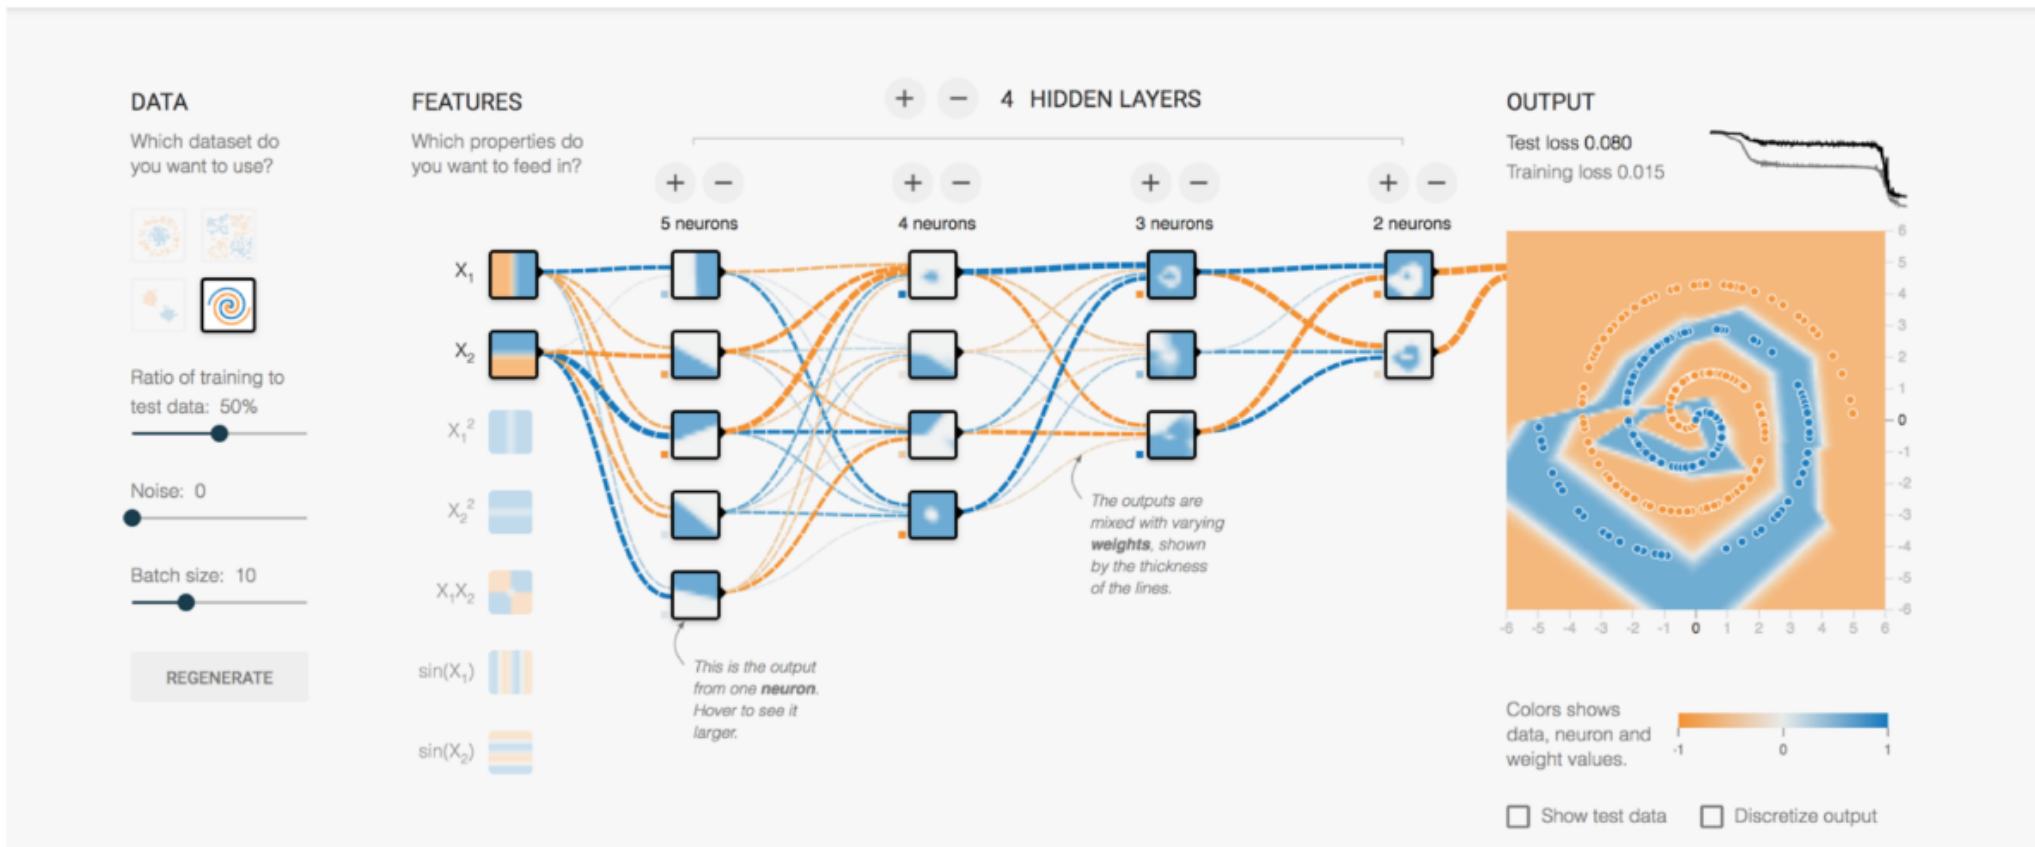
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

model = Sequential()

model.add(Conv2D(32, (3, 3), padding='same', input_shape=(32, 32, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu')) model.add(Dense(10))
model.add(Activation('softmax'))
```

Epoch
000,732Learning rate
0.03Activation
ReLURegularization
NoneRegularization rate
0Problem type
Classificationplayground.tensorflow.org

@poweredbyaltnet

Thank You!

@poweredbyaltnet

douglas@douglasstarnes.com

<http://douglasstarnes.com>

<https://github.com/douglasstarnes/indycodedeeplearning>